

```
In [ ]: import pandas as pd
import numpy as np
```

```
In [ ]: data = pd.read_csv("/Users/nishantsingh/Desktop/dataacuu/laptop_cleaned2.csv")
```

```
In [ ]: data.head()
```

Out[]:

	Unnamed: 0	Name	Brand	Price	Rating	Processor_brand	Processor_name	Processor_variant	Processor_gen	Core_per_processor	...	Graphics_name	Graphics_brand	Graphics_GB	Graphic
0	0	HP Victus 15-fb0157AX Gaming Laptop (AMD Ryzen...	HP	50399	4.30	AMD	AMD Ryzen 5	5600H	5.0	6.0	...	AMD Radeon RX 6500M	AMD	4.0	
1	1	Lenovo V15 G4 82YU00W7IN Laptop (AMD Ryzen 3 ...	Lenovo	26690	4.45	AMD	AMD Ryzen 3	7320U	7.0	4.0	...	AMD Radeon Graphics	AMD	NaN	
2	2	HP 15s-fq5007TU Laptop (12th Gen Core i3/ 8GB/...	HP	37012	4.65	Intel	Intel Core i3	1215U	12.0	6.0	...	Intel UHD Graphics	Intel	NaN	
3	3	Samsung Galaxy Book2 Pro 13 Laptop (12th Gen C...	Samsung	69990	4.75	Intel	Intel Core i5	1240P	12.0	12.0	...	Intel Iris Xe Graphics	Intel	NaN	
4	4	Tecno Megabook T1 Laptop (11th Gen Core i3/ 8G...	Tecno	23990	4.25	Intel	Intel Core i3	1115G4	11.0	2.0	...	Intel UHD Graphics	Intel	NaN	

5 rows × 29 columns

```
In [ ]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1020 entries, 0 to 1019
Data columns (total 29 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Unnamed: 0                            1020 non-null   int64
1   Name                                  1020 non-null   object
2   Brand                                 1020 non-null   object
3   Price                                 1020 non-null   int64
4   Rating                               1020 non-null   float64
5   Processor_brand                       1020 non-null   object
6   Processor_name                        1020 non-null   object
7   Processor_variant                     996 non-null    object
8   Processor_gen                         891 non-null    float64
9   Core_per_processor                    1008 non-null   float64
10  Total_processor                       573 non-null    float64
11  Execution_units                       573 non-null    float64
12  Low_Power_Cores                       1020 non-null   float64
13  Energy_Efficient_Units                1020 non-null   int64
14  Threads                               972 non-null    float64
15  RAM_GB                                1020 non-null   int64
16  RAM_type                              998 non-null    object
17  Storage_capacity_GB                   1020 non-null   int64
18  Storage_type                          1020 non-null   object
19  Graphics_name                         1018 non-null   object
20  Graphics_brand                        1018 non-null   object
21  Graphics_GB                           368 non-null    float64
22  Graphics_integreted                   1018 non-null   object
23  Display_size_inches                   1020 non-null   float64
24  Horizontal_pixel                       1020 non-null   int64
25  Vertical_pixel                         1020 non-null   int64
26  ppi                                    1020 non-null   float64
27  Touch_screen                          1020 non-null   bool
28  Operating_system                      1020 non-null   object
dtypes: bool(1), float64(10), int64(7), object(11)
memory usage: 224.2+ KB
```

```
In [ ]: unique_df = data.drop_duplicates()
        num_unique_rows = len(unique_df)
        print("Number of unique rows:", num_unique_rows)
```

Number of unique rows: 1020

```
In [ ]: data.describe()
```

Out []:

	Unnamed: 0	Price	Rating	Processor_gen	Core_per_processor	Total_processor	Execution_units	Low_Power_Cores	Energy_Efficient_Units	Threads	RAM_GB	Storage_capa
count	1020.000000	1020.000000	1020.000000	891.000000	1008.000000	573.000000	573.000000	1020.000000	1020.000000	972.000000	1020.000000	1020.
mean	509.500000	82063.474510	4.373676	10.450056	8.572421	3.926702	6.998255	0.086275	0.043137	12.817901	13.992157	627.
std	294.592939	66502.150607	0.233295	2.966579	4.375012	1.954429	2.680217	0.406531	0.203266	5.677459	7.189564	316.
min	0.000000	8000.000000	3.950000	1.000000	2.000000	1.000000	4.000000	0.000000	0.000000	2.000000	2.000000	32.
25%	254.750000	43990.000000	4.200000	7.000000	6.000000	2.000000	4.000000	0.000000	0.000000	8.000000	8.000000	512.
50%	509.500000	63689.500000	4.350000	12.000000	8.000000	4.000000	8.000000	0.000000	0.000000	12.000000	16.000000	512.
75%	764.250000	94990.000000	4.550000	13.000000	10.000000	6.000000	8.000000	0.000000	0.000000	16.000000	16.000000	512.
max	1019.000000	599990.000000	4.750000	14.000000	24.000000	12.000000	16.000000	2.000000	1.000000	32.000000	64.000000	4000.

In []:

```
categorical_data = data.select_dtypes(include=['object', 'category']).columns
categorical_data
```

Out []:

```
Index(['Name', 'Brand', 'Processor_brand', 'Processor_name',
       'Processor_variant', 'RAM_type', 'Storage_type', 'Graphics_name',
       'Graphics_brand', 'Graphics_integreted', 'Operating_system'],
      dtype='object')
```

In []:

```
continue_data = data.select_dtypes(include=['float64', 'int64']).columns
continue_data
```

Out []:

```
Index(['Unnamed: 0', 'Price', 'Rating', 'Processor_gen', 'Core_per_processor',
       'Total_processor', 'Execution_units', 'Low_Power_Cores',
       'Energy_Efficient_Units', 'Threads', 'RAM_GB', 'Storage_capacity_GB',
       'Graphics_GB', 'Display_size_inches', 'Horizontal_pixel',
       'Vertical_pixel', 'ppi'],
      dtype='object')
```

We now have to prep our data for our model Unnamed column was removed

we know that a laptops price will depend on the hardware features like ram,graphic, display etc. so we will select only that column

the columns will be

- 1 Brand 1020 non-null object 2 Price 1020 non-null int64
- 3 Rating 1020 non-null float64 5 Processor_name 1020 non-null object 8 Core_per_processor 1008 non-null float64 11 Low_Power_Cores 1020 non-null float64 12 Energy_Efficient_Units 1020 non-null int64
- 14 RAM_GB 1020 non-null int64
- 15 RAM_type 998 non-null object 16 Storage_capacity_GB 1020 non-null int64
- 17 Storage_type 1020 non-null object
- 19 Graphics_brand 1018 non-null object 21 Graphics_integreted 1018 non-null object 22 Display_size_inches 1020 non-null float64 23 Horizontal_pixel 1020 non-null int64
- 24 Vertical_pixel 1020 non-null int64
- 25 ppi 1020 non-null float64 26 Touch_screen 1020 non-null bool
- 27 Operating_system 1020 non-null object

these will be the columns which we will work on

```
In [ ]: selected_columns = [
    'Brand',
    'Price',
    'Rating',
    'Processor_name',
    'Core_per_processor',
    'Low_Power_Cores',
    'Energy_Efficient_Units',
    'RAM_GB',
    'RAM_type',
    'Storage_capacity_GB',
    'Storage_type',
    'Graphics_brand',
    'Graphics_integreted',
    'Display_size_inches',
    'Horizontal_pixel',
    'Vertical_pixel',
    'ppi',
    'Touch_screen',
    'Operating_system'
]
```

```
In [ ]: laptops = data[selected_columns]
laptops.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1020 entries, 0 to 1019
Data columns (total 19 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Brand                 1020 non-null  object
1   Price                 1020 non-null  int64
2   Rating                1020 non-null  float64
3   Processor_name        1020 non-null  object
4   Core_per_processor    1008 non-null  float64
5   Low_Power_Cores       1020 non-null  float64
6   Energy_Efficient_Units 1020 non-null  int64
7   RAM_GB                1020 non-null  int64
8   RAM_type              998 non-null   object
9   Storage_capacity_GB   1020 non-null  int64
10  Storage_type          1020 non-null  object
11  Graphics_brand         1018 non-null  object
12  Graphics_integreted    1018 non-null  object
13  Display_size_inches    1020 non-null  float64
14  Horizontal_pixel       1020 non-null  int64
15  Vertical_pixel         1020 non-null  int64
16  ppi                   1020 non-null  float64
17  Touch_screen           1020 non-null  bool
18  Operating_system       1020 non-null  object
dtypes: bool(1), float64(5), int64(6), object(7)
memory usage: 144.6+ KB
```

```
In [ ]: laptops.describe()
```

Out []:

	Price	Rating	Core_per_processor	Low_Power_Cores	Energy_Efficient_Units	RAM_GB	Storage_capacity_GB	Display_size_inches	Horizontal_pixel	Vertical_pixel	ppi
count	1020.000000	1020.000000	1008.000000	1020.000000	1020.000000	1020.000000	1020.000000	1020.000000	1020.000000	1020.000000	1020.000000
mean	82063.474510	4.373676	8.572421	0.086275	0.043137	13.992157	627.733333	15.163775	2035.512745	1214.019608	157.178265
std	66502.150607	0.233295	4.375012	0.406531	0.203266	7.189564	316.911679	1.001537	409.209289	306.863086	33.585713
min	8000.000000	3.950000	2.000000	0.000000	0.000000	2.000000	32.000000	11.600000	1080.000000	768.000000	100.450000
25%	43990.000000	4.200000	6.000000	0.000000	0.000000	8.000000	512.000000	14.000000	1920.000000	1080.000000	141.210000
50%	63689.500000	4.350000	8.000000	0.000000	0.000000	16.000000	512.000000	15.600000	1920.000000	1080.000000	141.210000
75%	94990.000000	4.550000	10.000000	0.000000	0.000000	16.000000	512.000000	15.600000	1920.000000	1200.000000	161.730000
max	599990.000000	4.750000	24.000000	2.000000	1.000000	64.000000	4000.000000	18.000000	3840.000000	2560.000000	337.930000

In []:

```
rows_with_null_ram_type = laptops[laptops['RAM_type'].isnull()]
rows_with_null_ram_type.shape
```

Out []: (22, 19)

In []:

```
laptops = laptops.dropna()
laptops.head()
```

Out []:

	Brand	Price	Rating	Processor_name	Core_per_processor	Low_Power_Cores	Energy_Efficient_Units	RAM_GB	RAM_type	Storage_capacity_GB	Storage_type	Graphics_brand	Graphics_integretec
0	HP	50399	4.30	AMD Ryzen 5	6.0	0.0	0	8	DDR4	512	SSD	AMD	False
1	Lenovo	26690	4.45	AMD Ryzen 3	4.0	0.0	0	8	LPDDR5	512	SSD	AMD	False
2	HP	37012	4.65	Intel Core i3	6.0	0.0	0	8	DDR4	512	SSD	Intel	False
3	Samsung	69990	4.75	Intel Core i5	12.0	0.0	0	16	LPDDR5	512	SSD	Intel	False
4	Tecno	23990	4.25	Intel Core i3	2.0	0.0	0	8	LPDDR4	512	SSD	Intel	False

Getting a visualization of our data

In []:

```
import matplotlib.pyplot as plt
```

In []:

```
# let us understand how different features lead to the price our laptops
```

'Brand', 'Price', 'Rating', 'Processor_name', 'Core_per_processor', 'Low_Power_Cores', 'Energy_Efficient_Units', 'RAM_GB', 'RAM_type', 'Storage_capacity_GB', 'Storage_type', 'Graphics_brand', 'Graphics_integrated', 'Display_size_inches', 'Horizontal_pixel', 'Vertical_pixel', 'ppi', 'Touch_screen', 'Operating_system'


lowcore

energy effi graphic touch

In []:

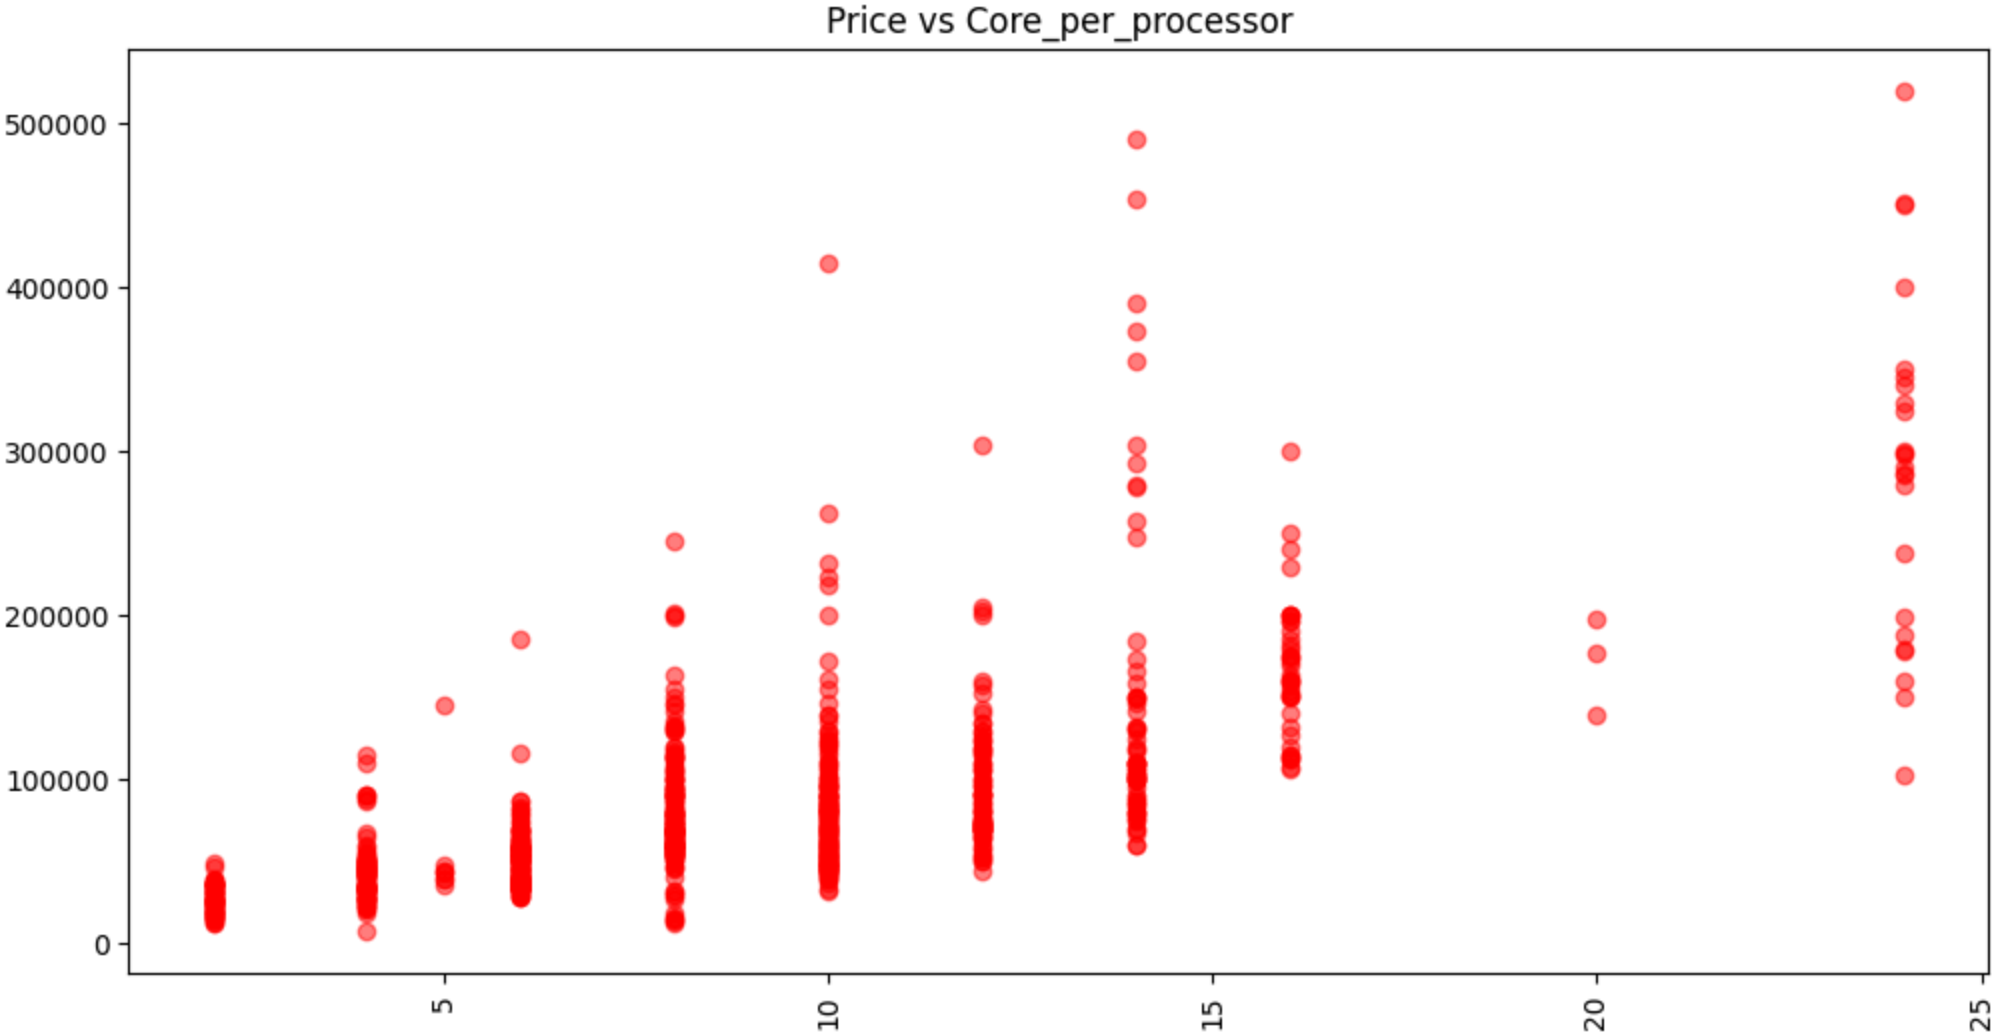
```
no_scatter_plot = ["Low_Power_Cores", "Energy_Efficient_Units", "Graphics_integretec", "Touch_screen"]
for cols in selected_columns:
```

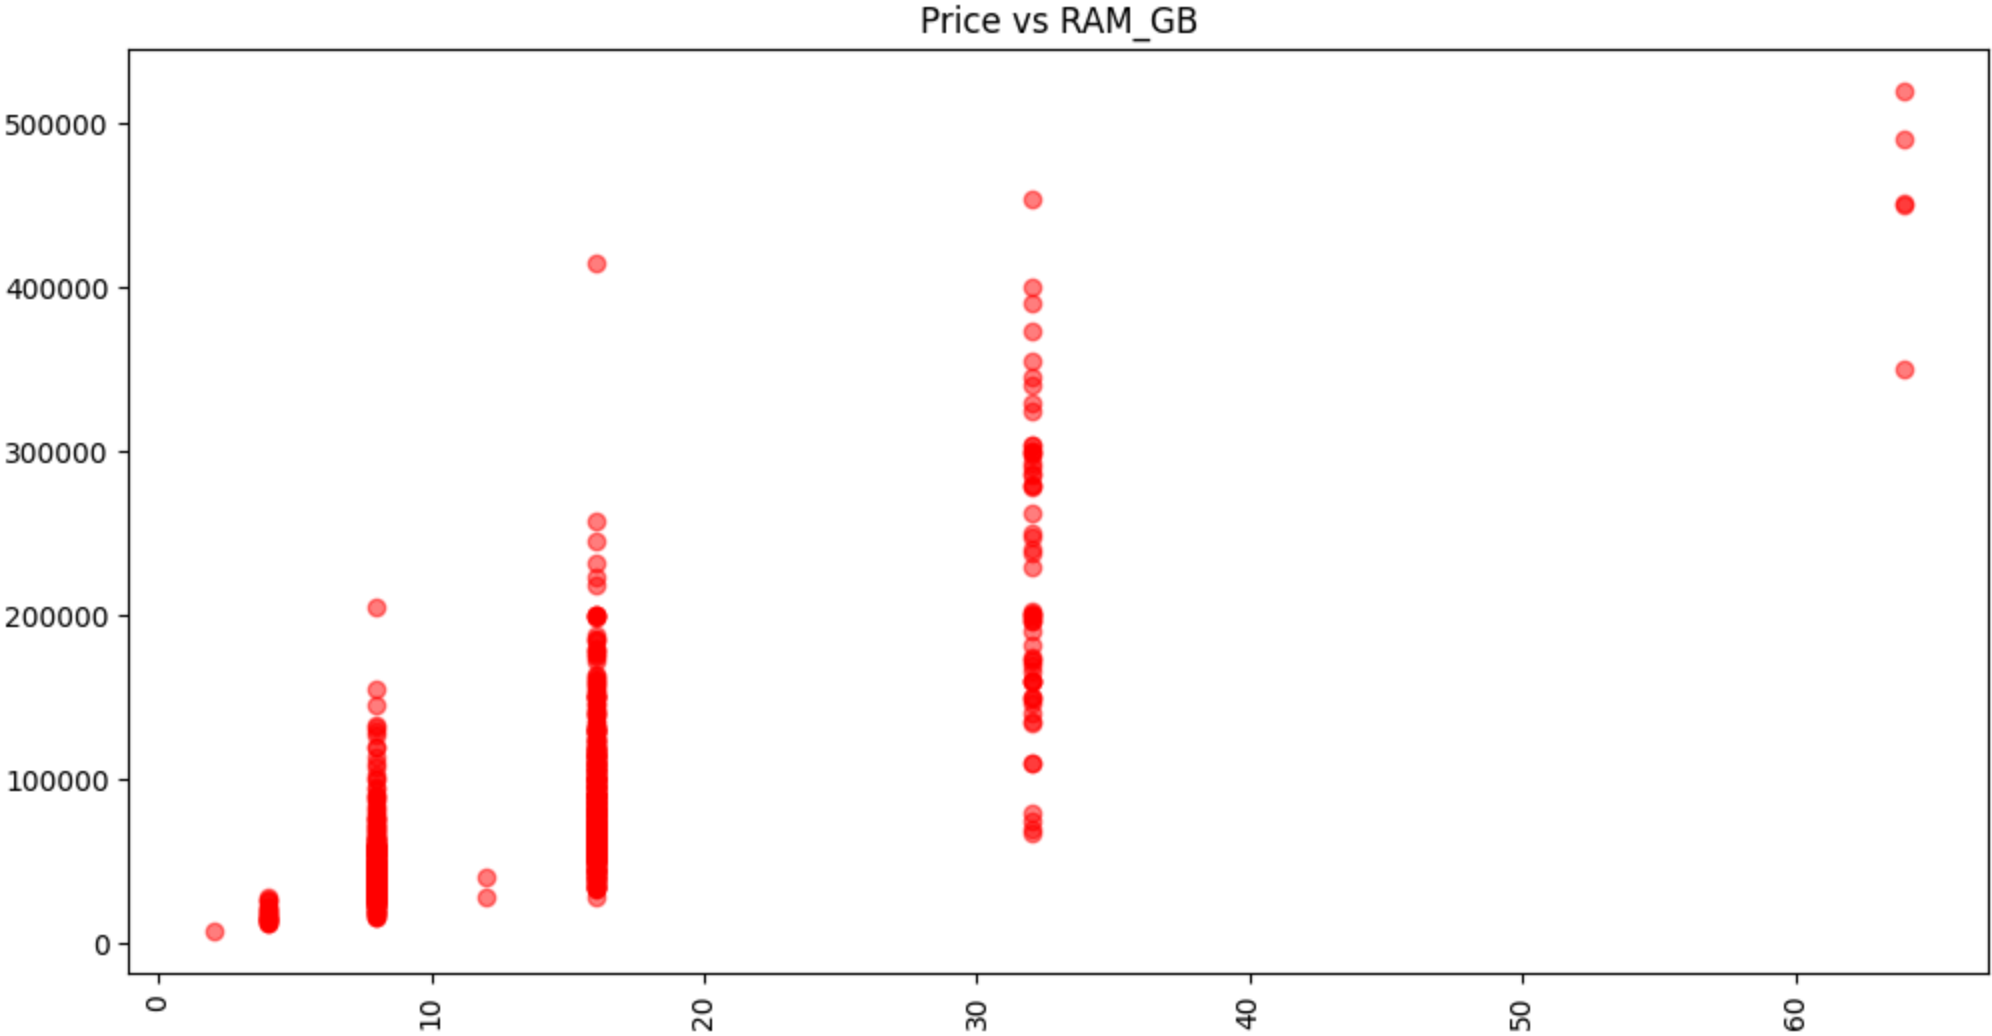
```
if (cols not in no_scatter_plot) & (cols != "Price"):
    plt.figure(figsize=(12, 6))
    plt.scatter(laptops[cols], laptops["Price"],alpha=0.5,color='red')
    plt.xticks(rotation=90)
    plt.title(f"Price vs {cols}")
    plt.show()
```

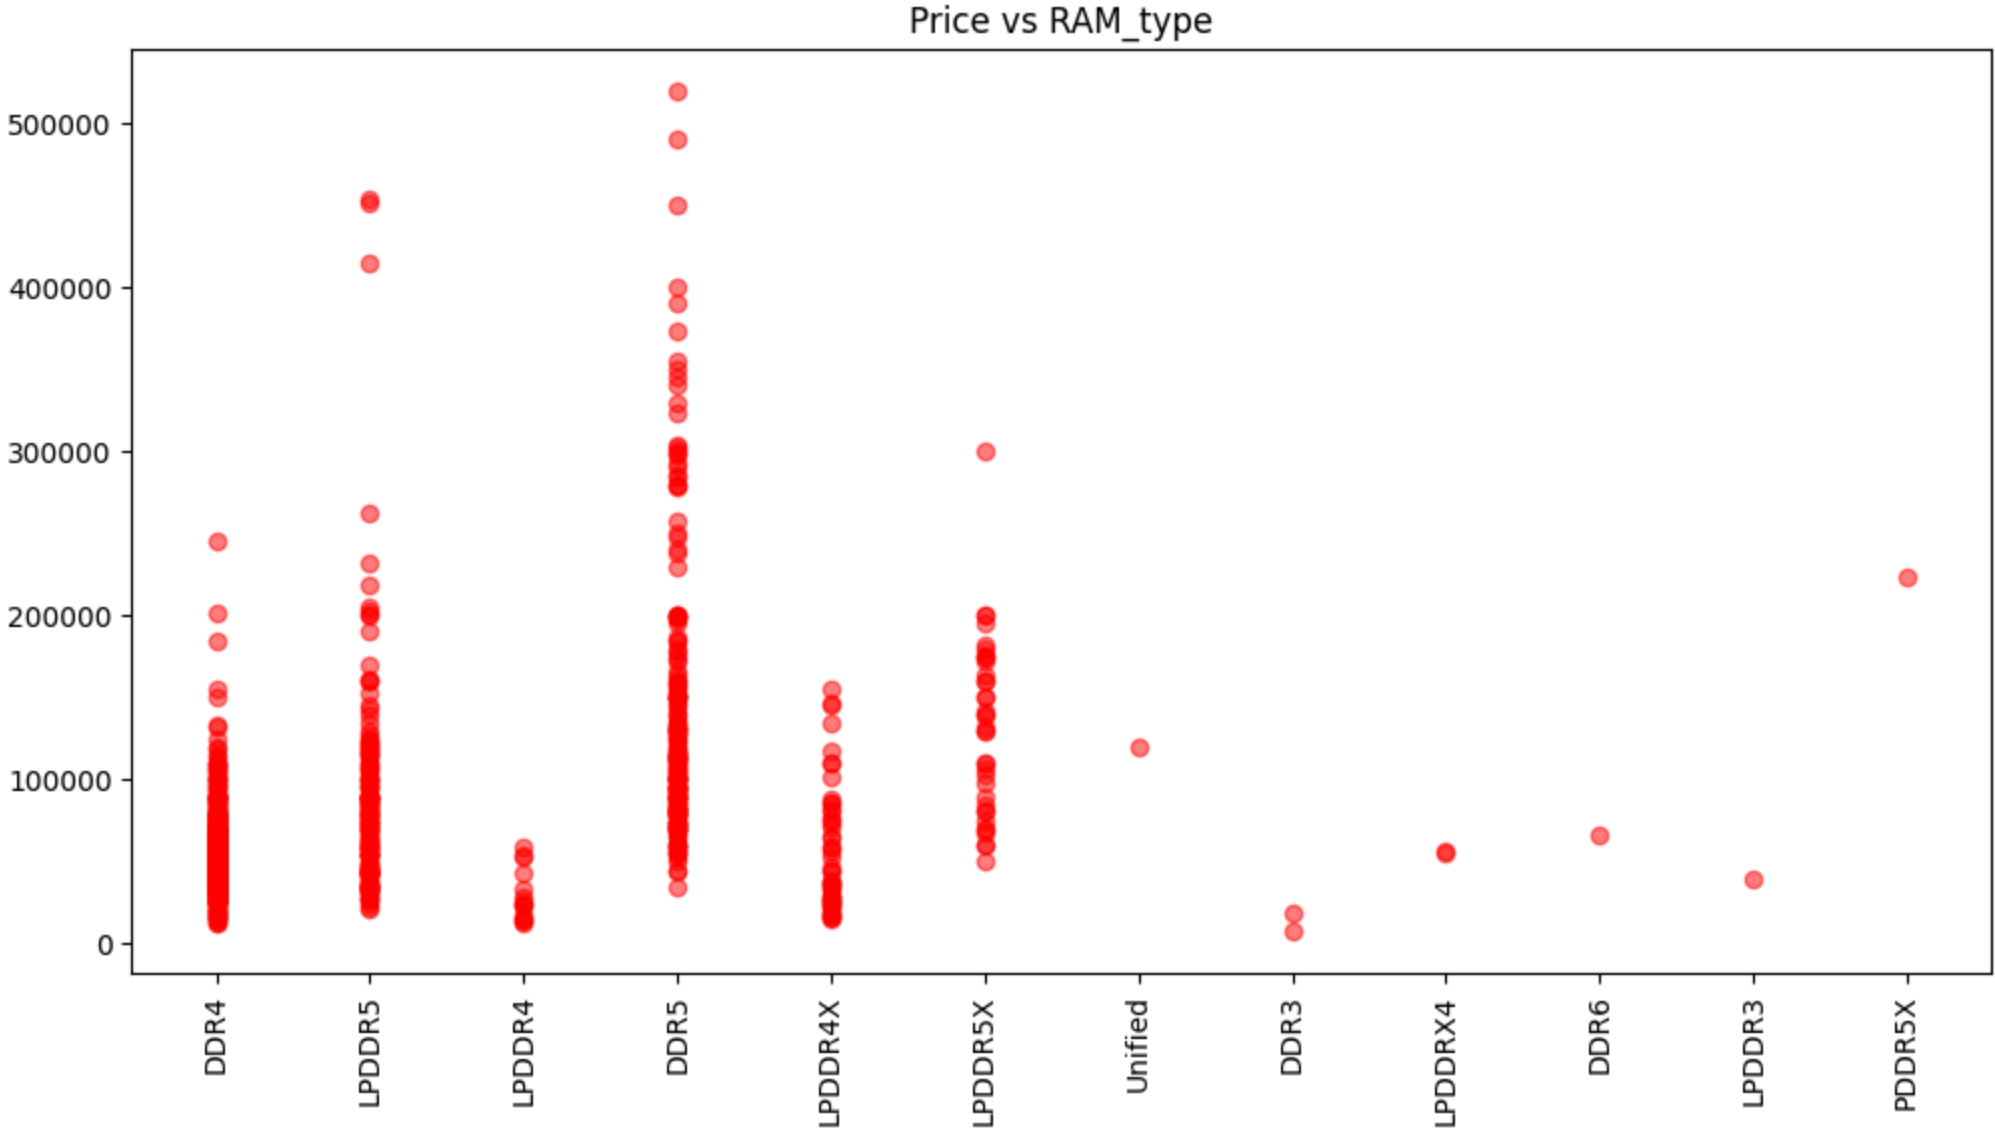





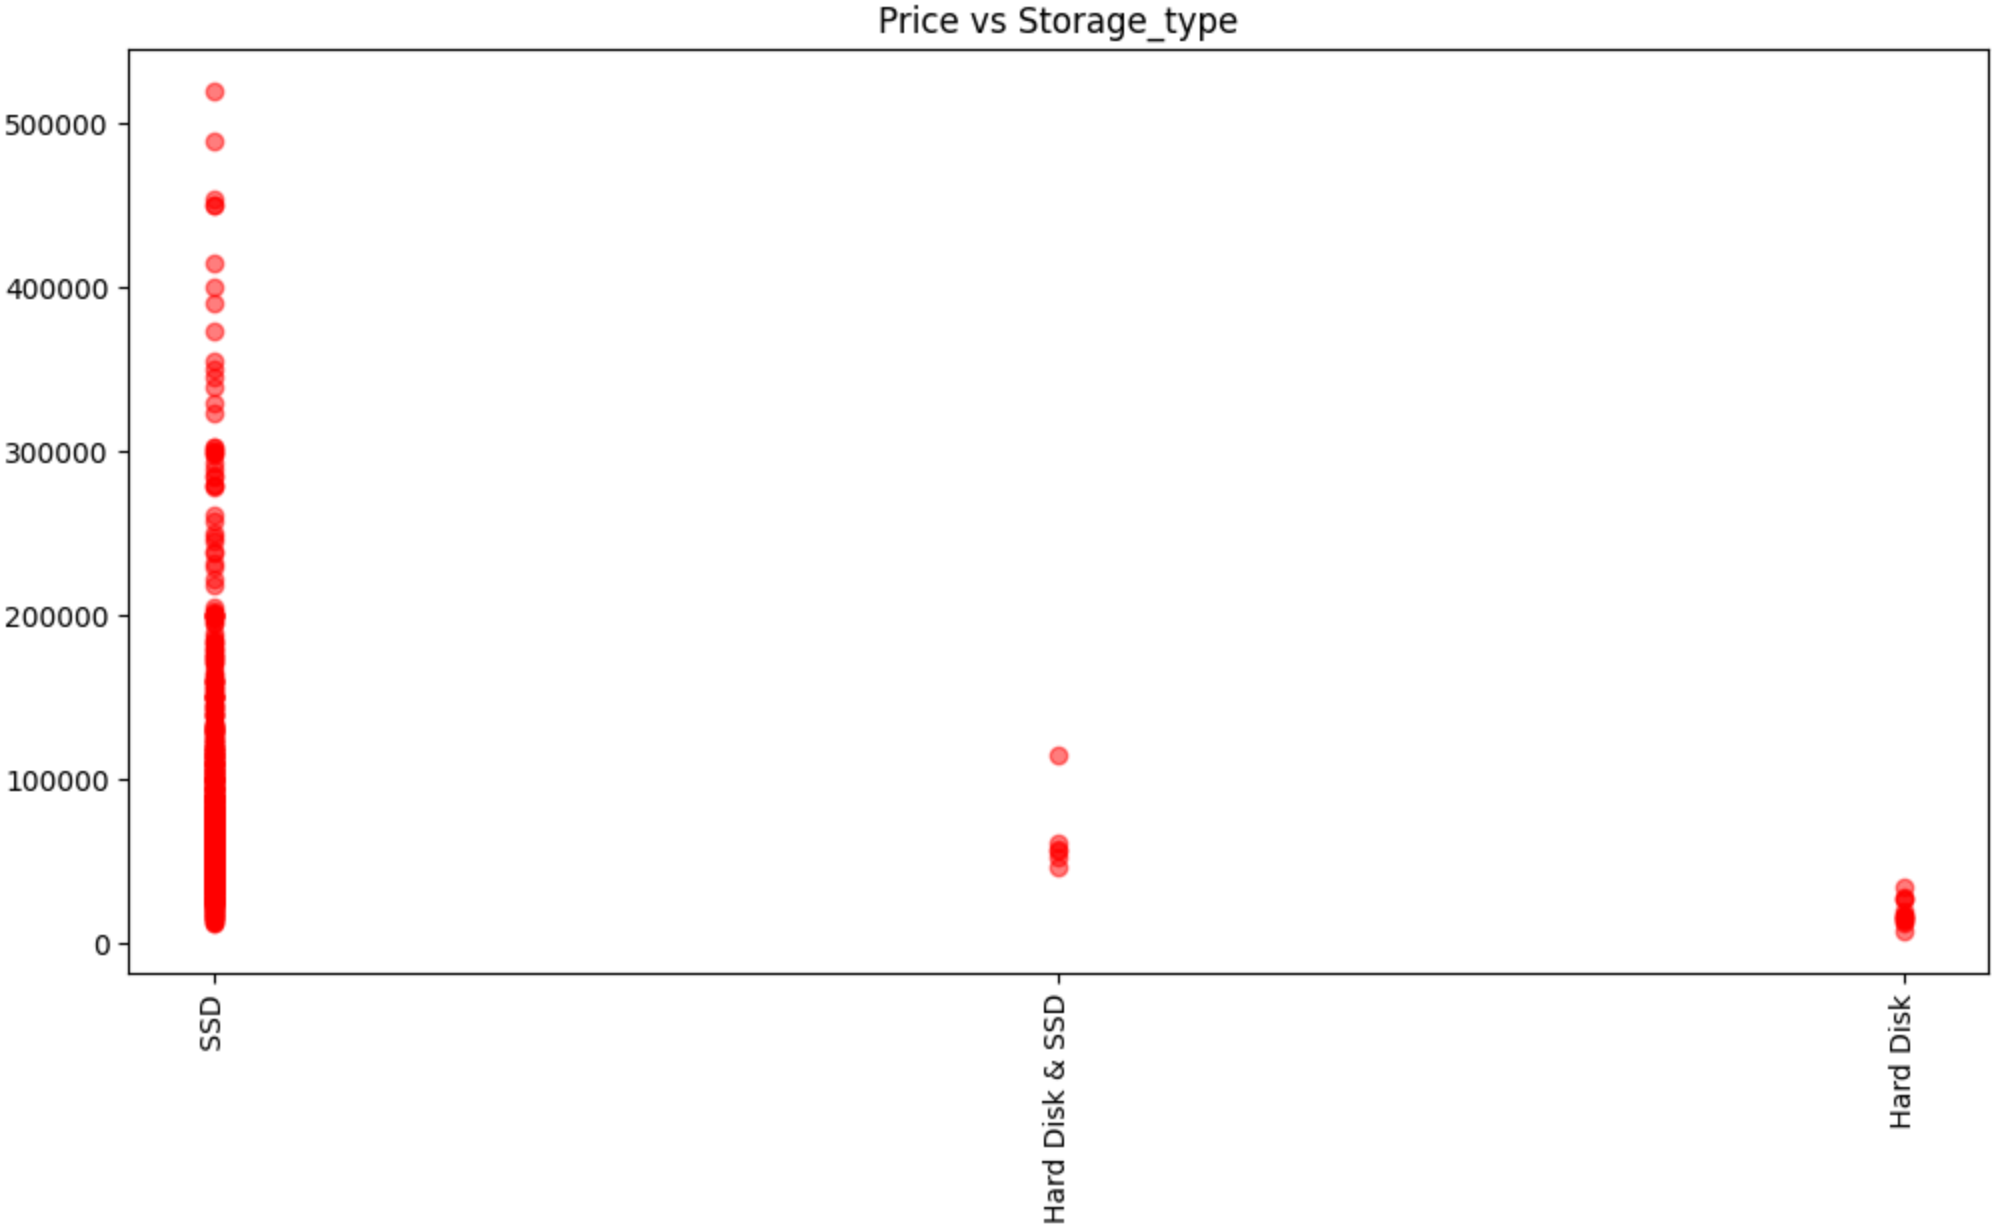


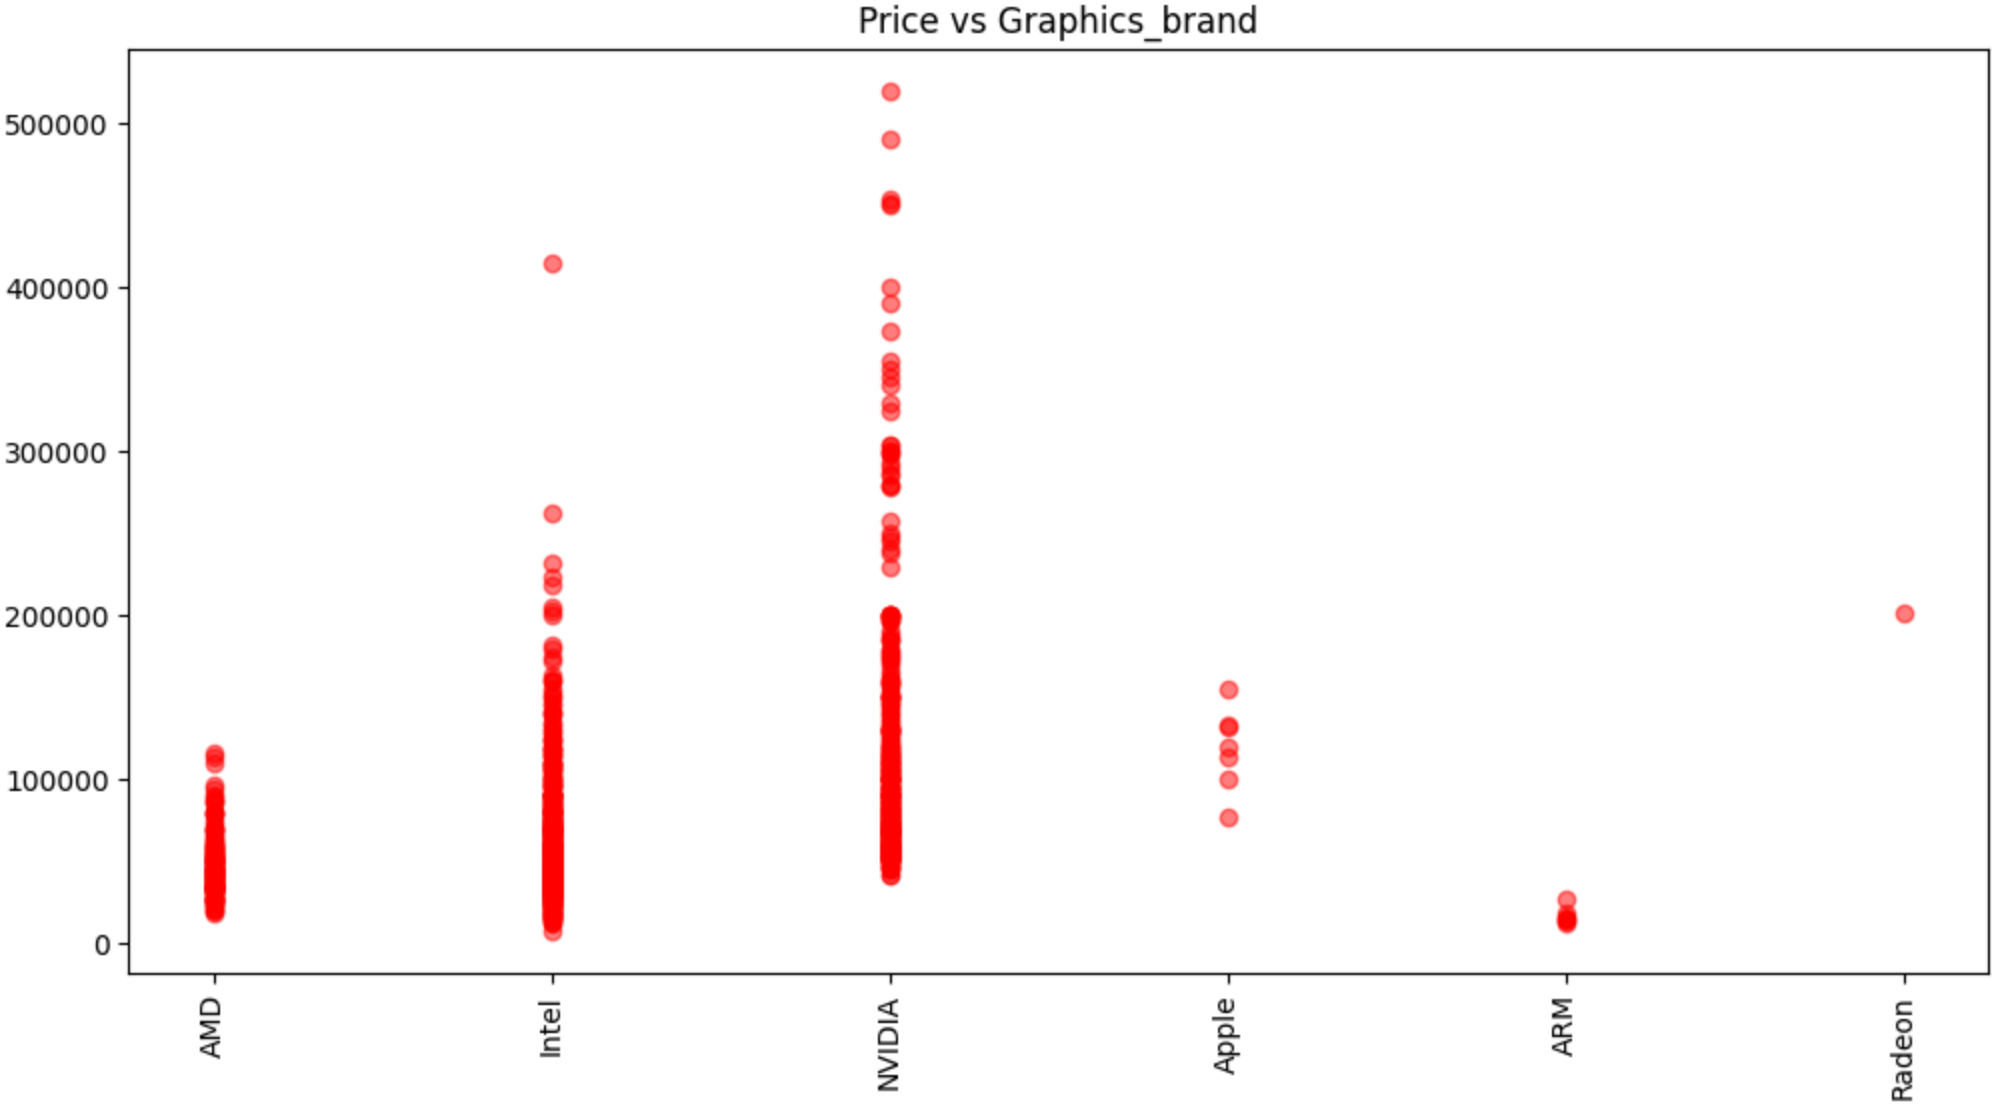














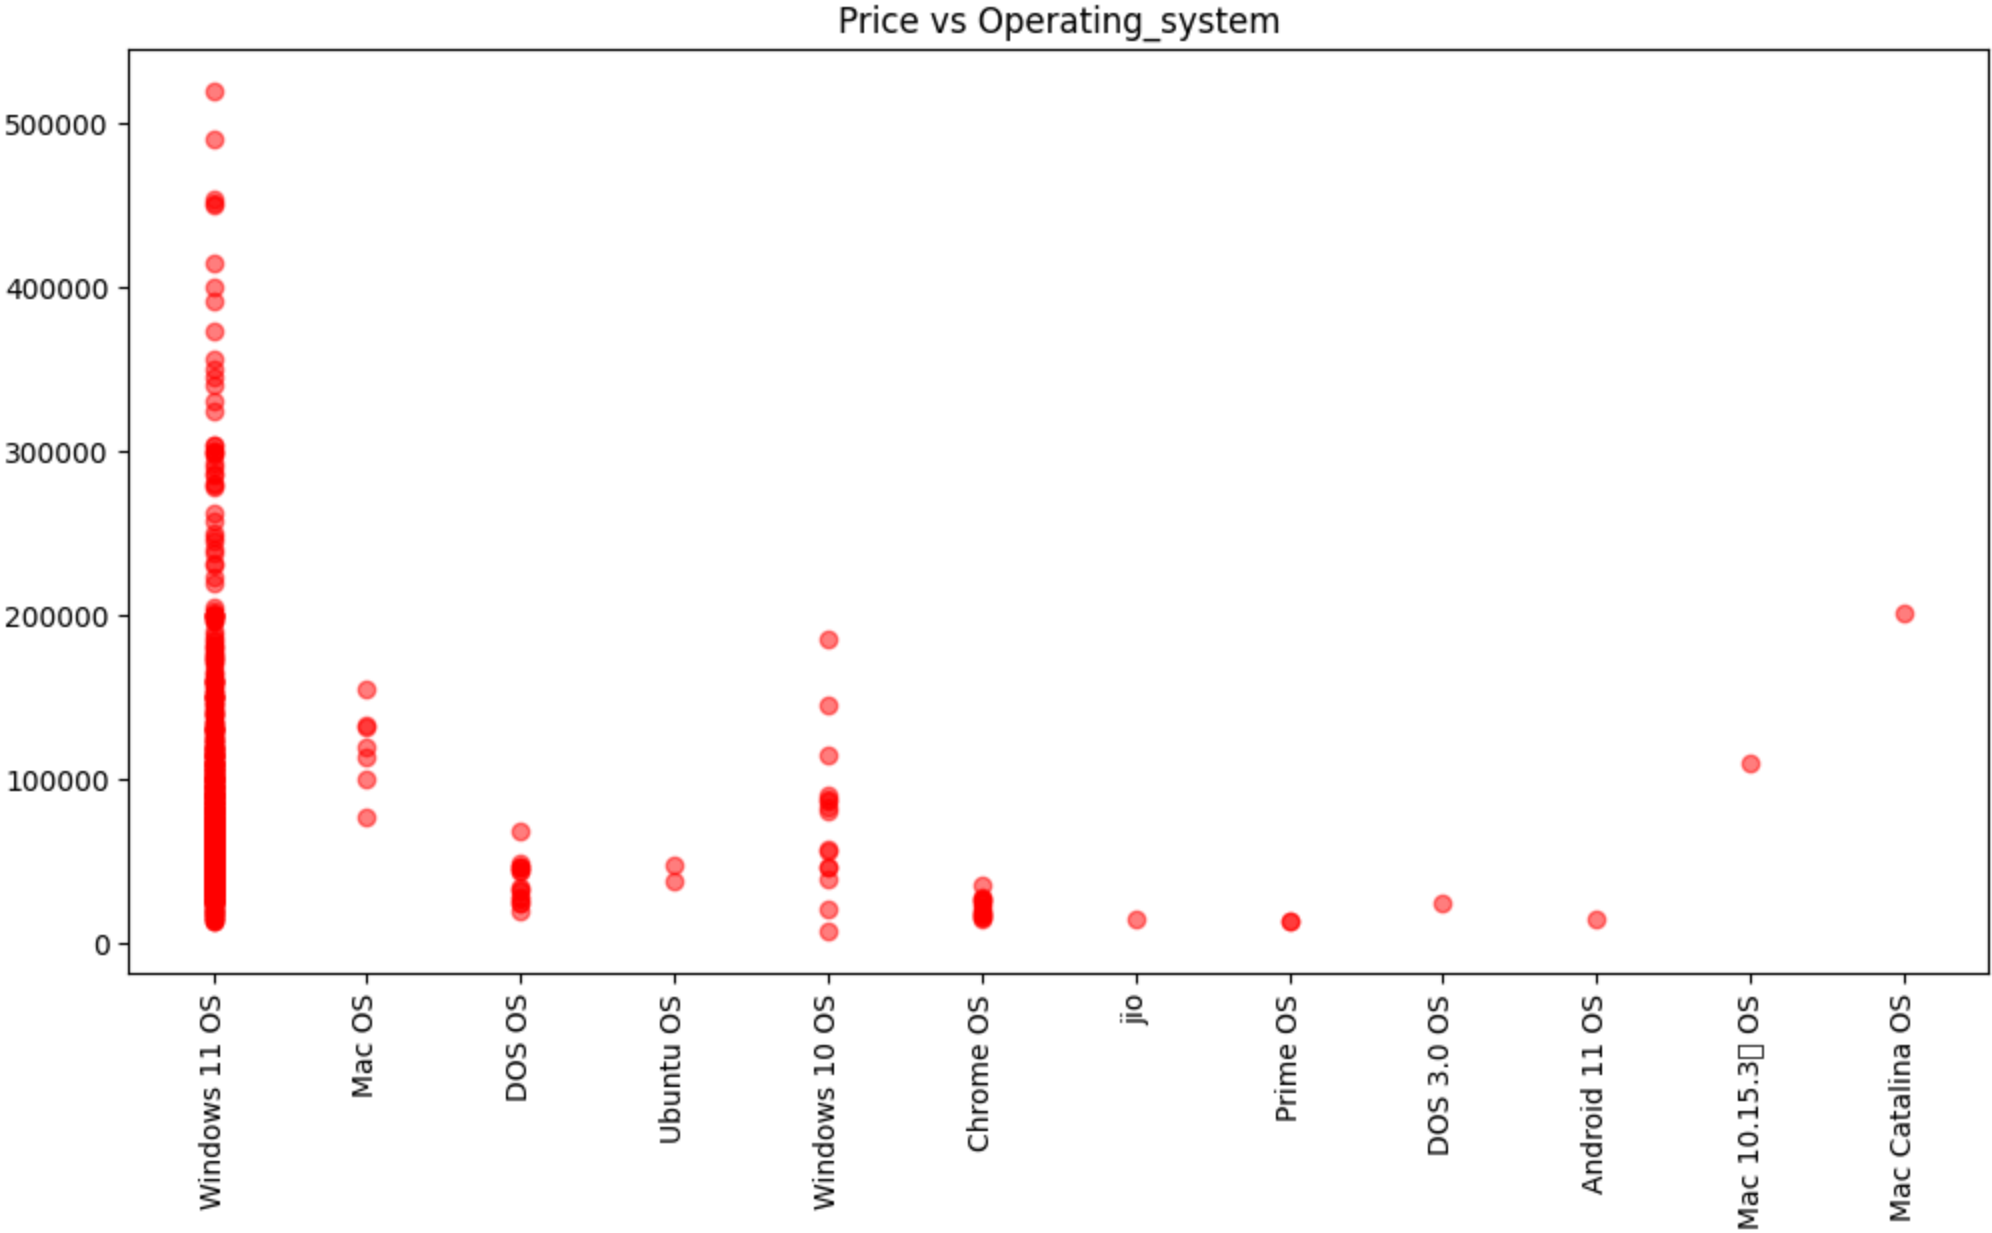








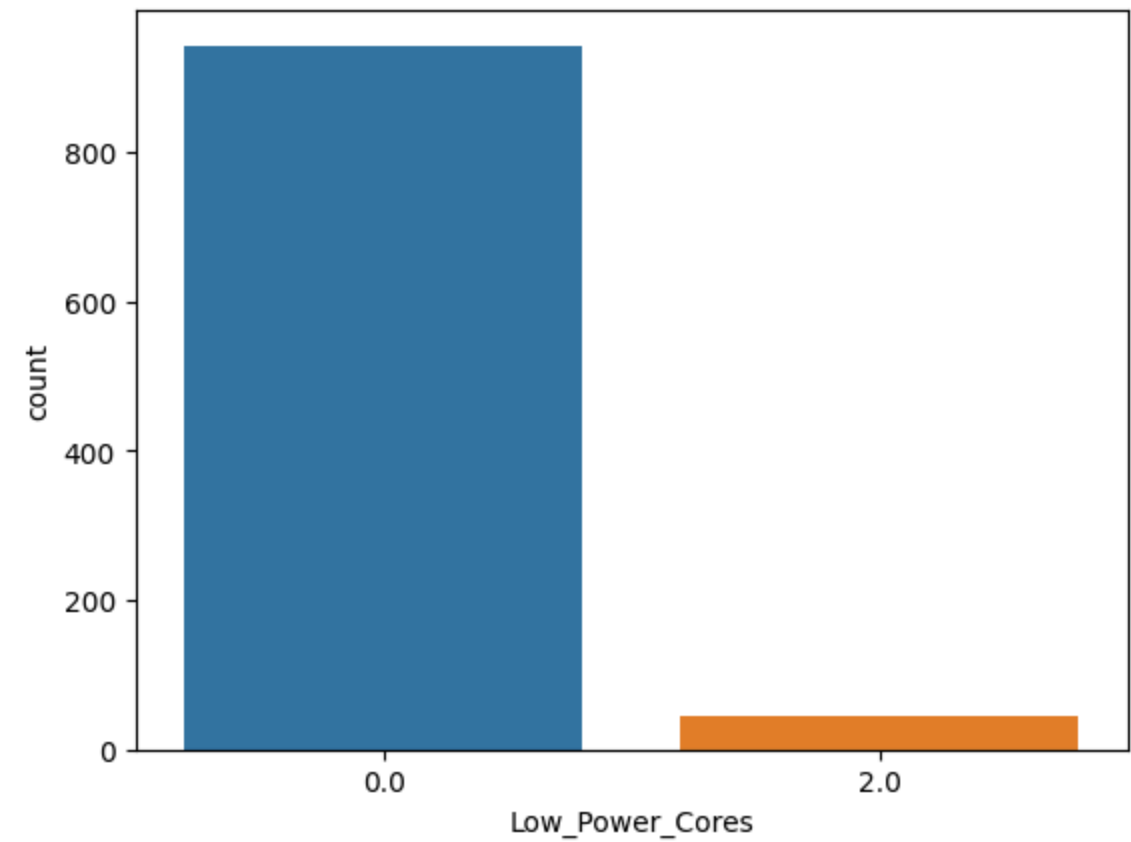
```
/Users/nishantsingh/Library/Python/3.11/lib/python/site-packages/IPython/core/pylabtools.py:152: UserWarning: Glyph 9 ( ) missing from current font.  
fig.canvas.print_figure(bytes_io, **kw)
```



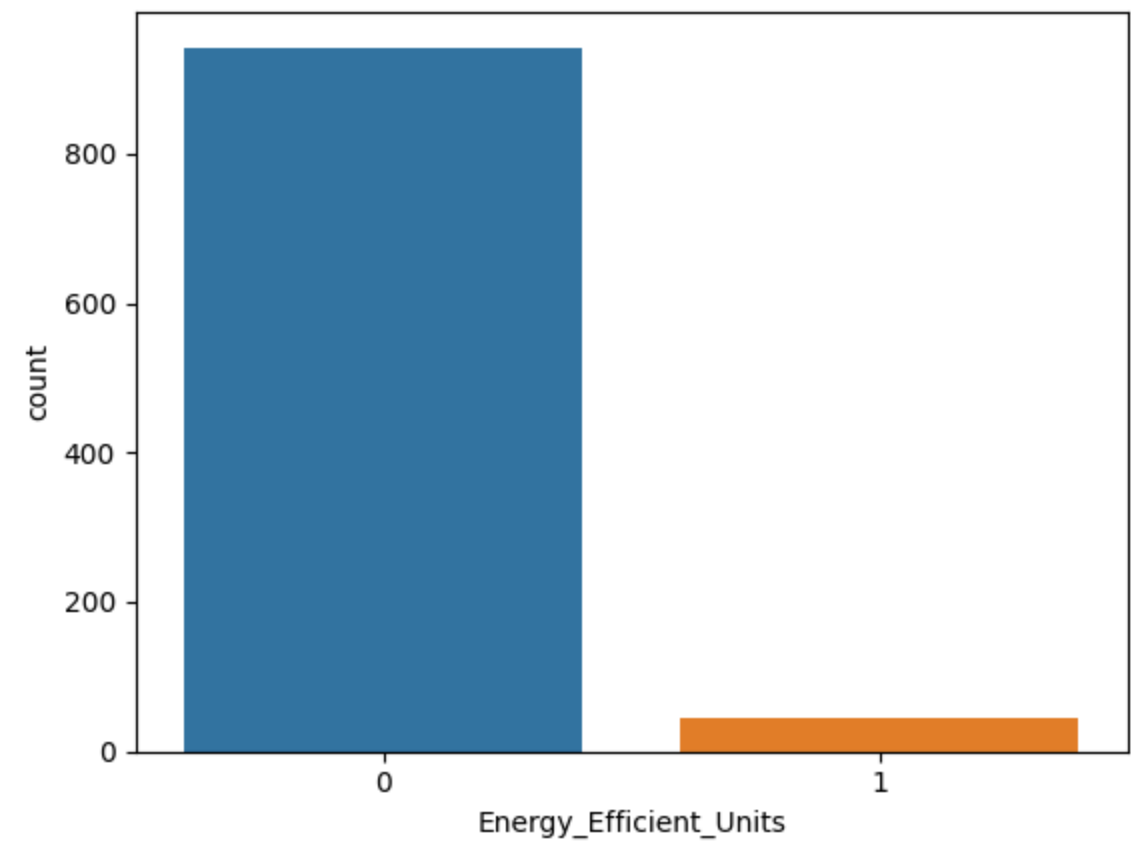
```
In [ ]: bar_plot = ["Low_Power_Cores","Energy_Efficient_Units","Graphics_integreted","Touch_screen"]

import seaborn as sns
for cols in bar_plot:
    sns.countplot(x=cols, data=laptops)
plt.show()
```

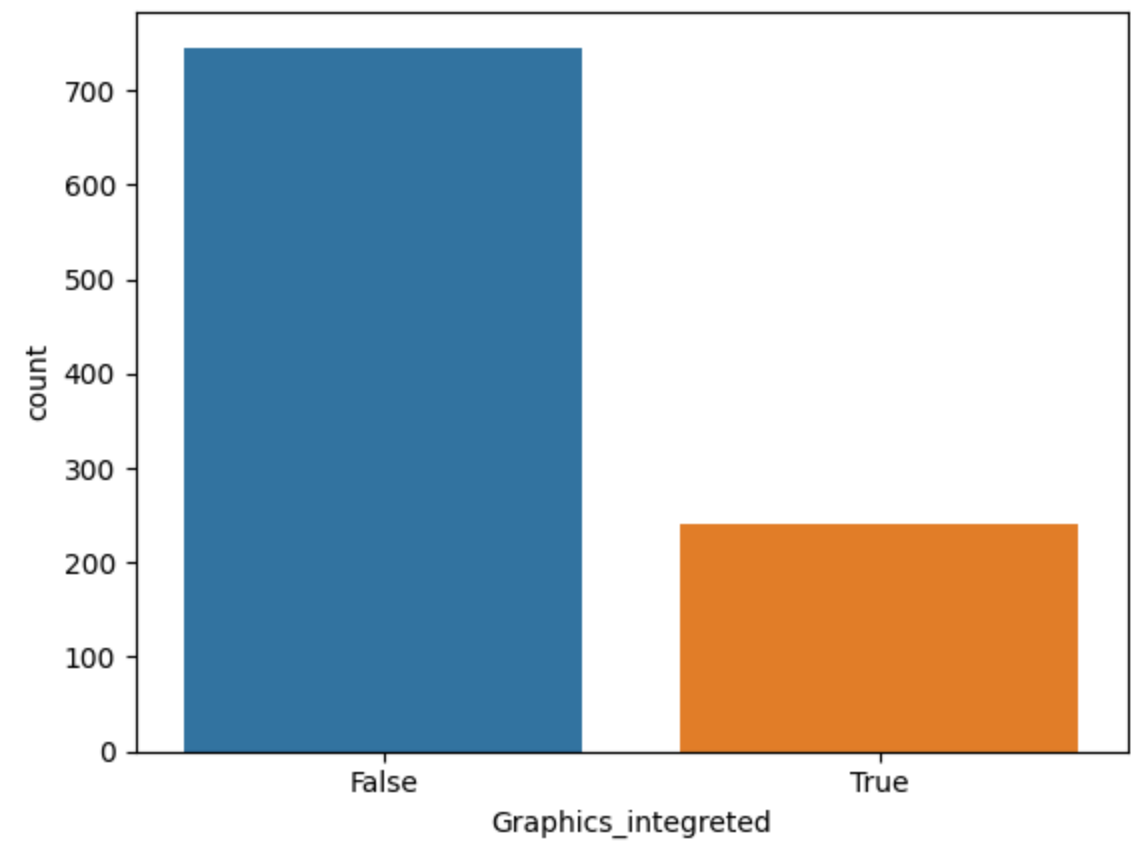
/Library/Frameworks/Python.framework/Versions/3.11/lib/python3.11/site-packages/seaborn/_oldcore.py:1498: FutureWarning: is_categorical_dtype is deprecated and will be removed in a future version. Use isinstance(dtype, CategoricalDtype) instead
if pd.api.types.is_categorical_dtype(vector):
/Library/Frameworks/Python.framework/Versions/3.11/lib/python3.11/site-packages/seaborn/_oldcore.py:1498: FutureWarning: is_categorical_dtype is deprecated and will be removed in a future version. Use isinstance(dtype, CategoricalDtype) instead
if pd.api.types.is_categorical_dtype(vector):
/Library/Frameworks/Python.framework/Versions/3.11/lib/python3.11/site-packages/seaborn/_oldcore.py:1498: FutureWarning: is_categorical_dtype is deprecated and will be removed in a future version. Use isinstance(dtype, CategoricalDtype) instead
if pd.api.types.is_categorical_dtype(vector):



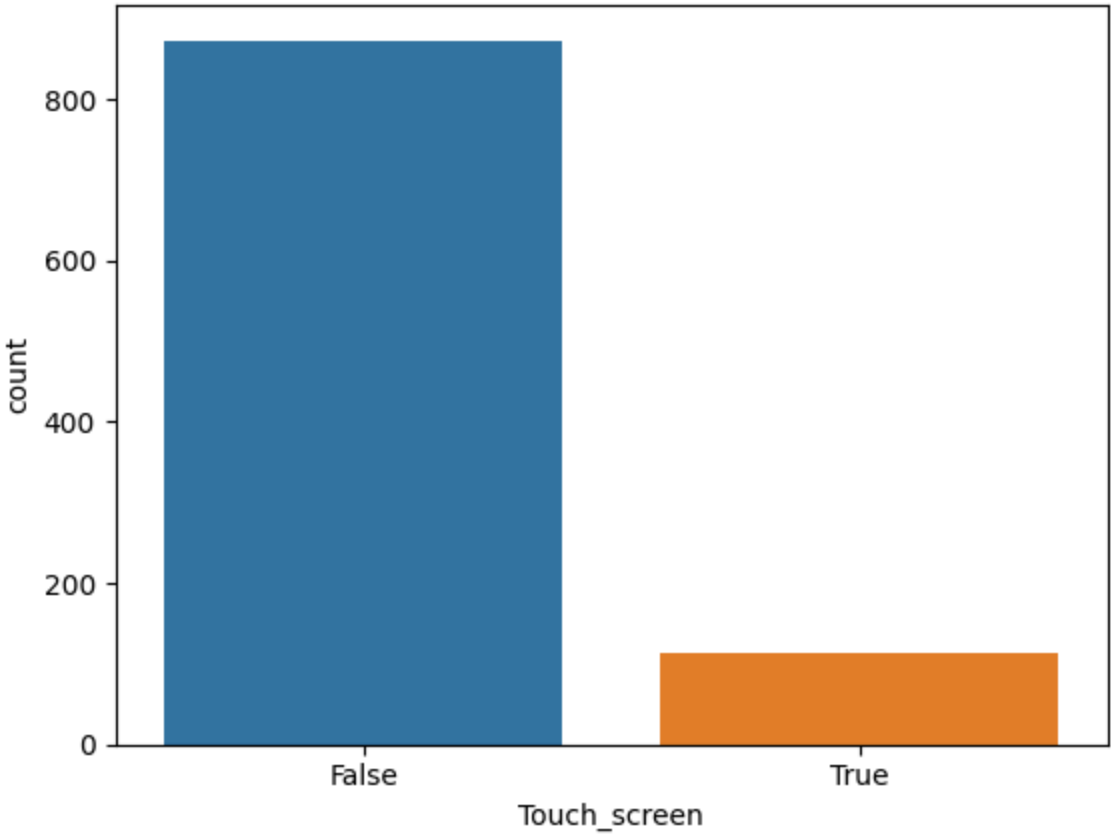
```
/Library/Frameworks/Python.framework/Versions/3.11/lib/python3.11/site-packages/seaborn/_oldcore.py:1498: FutureWarning: is_categorical_dtype is deprecated and will be removed in a future version. Use isinstance(dtype, CategoricalDtype) instead
  if pd.api.types.is_categorical_dtype(vector):
/Library/Frameworks/Python.framework/Versions/3.11/lib/python3.11/site-packages/seaborn/_oldcore.py:1498: FutureWarning: is_categorical_dtype is deprecated and will be removed in a future version. Use isinstance(dtype, CategoricalDtype) instead
  if pd.api.types.is_categorical_dtype(vector):
/Library/Frameworks/Python.framework/Versions/3.11/lib/python3.11/site-packages/seaborn/_oldcore.py:1498: FutureWarning: is_categorical_dtype is deprecated and will be removed in a future version. Use isinstance(dtype, CategoricalDtype) instead
  if pd.api.types.is_categorical_dtype(vector):
```



```
/Library/Frameworks/Python.framework/Versions/3.11/lib/python3.11/site-packages/seaborn/_oldcore.py:1498: FutureWarning: is_categorical_dtype is deprecated and will be removed in a future version. Use isinstance(dtype, CategoricalDtype) instead
    if pd.api.types.is_categorical_dtype(vector):
/Library/Frameworks/Python.framework/Versions/3.11/lib/python3.11/site-packages/seaborn/_oldcore.py:1498: FutureWarning: is_categorical_dtype is deprecated and will be removed in a future version. Use isinstance(dtype, CategoricalDtype) instead
    if pd.api.types.is_categorical_dtype(vector):
/Library/Frameworks/Python.framework/Versions/3.11/lib/python3.11/site-packages/seaborn/_oldcore.py:1498: FutureWarning: is_categorical_dtype is deprecated and will be removed in a future version. Use isinstance(dtype, CategoricalDtype) instead
    if pd.api.types.is_categorical_dtype(vector):
```



```
/Library/Frameworks/Python.framework/Versions/3.11/lib/python3.11/site-packages/seaborn/_oldcore.py:1498: FutureWarning: is_categorical_dtype is deprecated and will be removed in a future version. Use isinstance(dtype, CategoricalDtype) instead
  if pd.api.types.is_categorical_dtype(vector):
/Library/Frameworks/Python.framework/Versions/3.11/lib/python3.11/site-packages/seaborn/_oldcore.py:1498: FutureWarning: is_categorical_dtype is deprecated and will be removed in a future version. Use isinstance(dtype, CategoricalDtype) instead
  if pd.api.types.is_categorical_dtype(vector):
/Library/Frameworks/Python.framework/Versions/3.11/lib/python3.11/site-packages/seaborn/_oldcore.py:1498: FutureWarning: is_categorical_dtype is deprecated and will be removed in a future version. Use isinstance(dtype, CategoricalDtype) instead
  if pd.api.types.is_categorical_dtype(vector):
```



```
In [ ]: selected_columns_2 = [  
    'Price',  
    'Rating',  
    'Core_per_processor',  
    'Low_Power_Cores',  
    'Energy_Efficient_Units',  
    'RAM_GB',  
    'Storage_capacity_GB',  
    'Storage_type',  
    'Graphics_integreted',  
    'Display_size_inches',  
    'Horizontal_pixel',  
    'Vertical_pixel',  
    'ppi',  
    'Touch_screen',  
    'Operating_system'  
]  
  
In [ ]: df = laptops[selected_columns_2]  
df.head()
```

Out []:

	Price	Rating	Core_per_processor	Low_Power_Cores	Energy_Efficient_Units	RAM_GB	Storage_capacity_GB	Storage_type	Graphics_integreted	Display_size_inches	Horizontal_pixel	Vertical_pixel
0	50399	4.30	6.0	0.0	0	8	512	SSD	False	15.6	1920	1080
1	26690	4.45	4.0	0.0	0	8	512	SSD	False	15.6	1920	1080
2	37012	4.65	6.0	0.0	0	8	512	SSD	False	15.6	1920	1080
3	69990	4.75	12.0	0.0	0	16	512	SSD	False	13.3	1080	1920
4	23990	4.25	2.0	0.0	0	8	512	SSD	False	15.6	1920	1080

In []:

```
from sklearn.preprocessing import LabelEncoder
label_encoder = LabelEncoder()

encoding_labels = ['Storage_type', 'Operating_system']

for cols in encoding_labels:
    df[cols] = label_encoder.fit_transform(df[cols])

df['Graphics_integreted'] = df['Graphics_integreted'].astype(int)
df['Touch_screen'] = df['Touch_screen'].astype(int)
```

/var/folders/kn/x036pv7d2sgg8wy4n1b8614m0000gn/T/ipykernel_11764/2353218521.py:7: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
df[cols] = label_encoder.fit_transform(df[cols])

/var/folders/kn/x036pv7d2sgg8wy4n1b8614m0000gn/T/ipykernel_11764/2353218521.py:7: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
df[cols] = label_encoder.fit_transform(df[cols])

/var/folders/kn/x036pv7d2sgg8wy4n1b8614m0000gn/T/ipykernel_11764/2353218521.py:9: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
df['Graphics_integreted'] = df['Graphics_integreted'].astype(int)

/var/folders/kn/x036pv7d2sgg8wy4n1b8614m0000gn/T/ipykernel_11764/2353218521.py:10: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
df['Touch_screen'] = df['Touch_screen'].astype(int)

In []:


```
df.corr()
```


Out []:

	Price	Rating	Core_per_processor	Low_Power_Cores	Energy_Efficient_Units	RAM_GB	Storage_capacity_GB	Storage_type	Graphics_integreted	Display_size_inches	Horizontal_pixel
Price	1.000000	-0.041996	0.719818	0.219841	0.219841	0.774043	0.671881	0.086042	-0.234395	0.242106	0.651431
Rating	-0.041996	1.000000	-0.004586	-0.040597	-0.040597	-0.042779	-0.028134	0.042342	0.016629	-0.036497	0.007940
Core_per_processor	0.719818	-0.004586	1.000000	0.314600	0.314600	0.641045	0.532406	0.036563	-0.132618	0.192844	0.524051
Low_Power_Cores	0.219841	-0.040597	0.314600	1.000000	1.000000	0.249824	0.202761	0.012319	-0.031354	-0.061839	0.320568
Energy_Efficient_Units	0.219841	-0.040597	0.314600	1.000000	1.000000	0.249824	0.202761	0.012319	-0.031354	-0.061839	0.320568
RAM_GB	0.774043	-0.042779	0.641045	0.249824	0.249824	1.000000	0.671089	0.130616	-0.133784	0.222448	0.558447
Storage_capacity_GB	0.671881	-0.028134	0.532406	0.202761	0.202761	0.671089	1.000000	0.264789	-0.150166	0.272962	0.468446
Storage_type	0.086042	0.042342	0.036563	0.012319	0.012319	0.130616	0.264789	1.000000	-0.033963	0.226965	0.126722
Graphics_integreted	-0.234395	0.016629	-0.132618	-0.031354	-0.031354	-0.133784	-0.150166	-0.033963	1.000000	-0.196917	-0.130271
Display_size_inches	0.242106	-0.036497	0.192844	-0.061839	-0.061839	0.222448	0.272962	0.226965	-0.196917	1.000000	0.128469
Horizontal_pixel	0.651431	-0.046182	0.524051	0.320568	0.320568	0.558447	0.468446	0.126722	-0.130271	0.128469	1.000000
Vertical_pixel	0.656980	0.007940	0.511406	0.303999	0.303999	0.519890	0.426760	0.100602	-0.146771	0.040142	0.000000
ppi	0.571938	-0.011371	0.459006	0.343733	0.343733	0.468720	0.362385	0.058470	-0.073102	-0.239137	0.000000
Touch_screen	0.210706	-0.021116	0.118444	0.215280	0.215280	0.105867	0.042563	-0.046663	0.010431	-0.264039	0.000000
Operating_system	0.126446	-0.026248	0.164709	0.046271	0.046271	0.205819	0.202309	0.456387	-0.090805	0.172370	0.000000

In []:

```
correlation_matrix = df.corr()
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f", linewidths=0.5)
plt.title('Correlation Matrix')
plt.show()
```



```
In [ ]: columns=['Rating', 'Core_per_processor', 'RAM_GB', 'Storage_capacity_GB',  
              'Graphics_integrated', 'Display_size_inches', 'Horizontal_pixel',  
              'Vertical_pixel', 'ppi', 'Touch_screen']  
  
In [ ]: from sklearn.preprocessing import StandardScaler
```

```
# Create a StandardScaler object
scaler = StandardScaler()

# Fit and transform your data
scaled_data = scaler.fit_transform(df)
```

```
In [ ]: from sklearn.model_selection import train_test_split
        from sklearn.metrics import mean_squared_error
        from math import sqrt
        from sklearn.metrics import r2_score

        y = df.iloc[:, 0]
        X = df.iloc[:, 1:]
        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [ ]: from sklearn.ensemble import RandomForestRegressor

        model = RandomForestRegressor()
        model.fit(X_train, y_train)

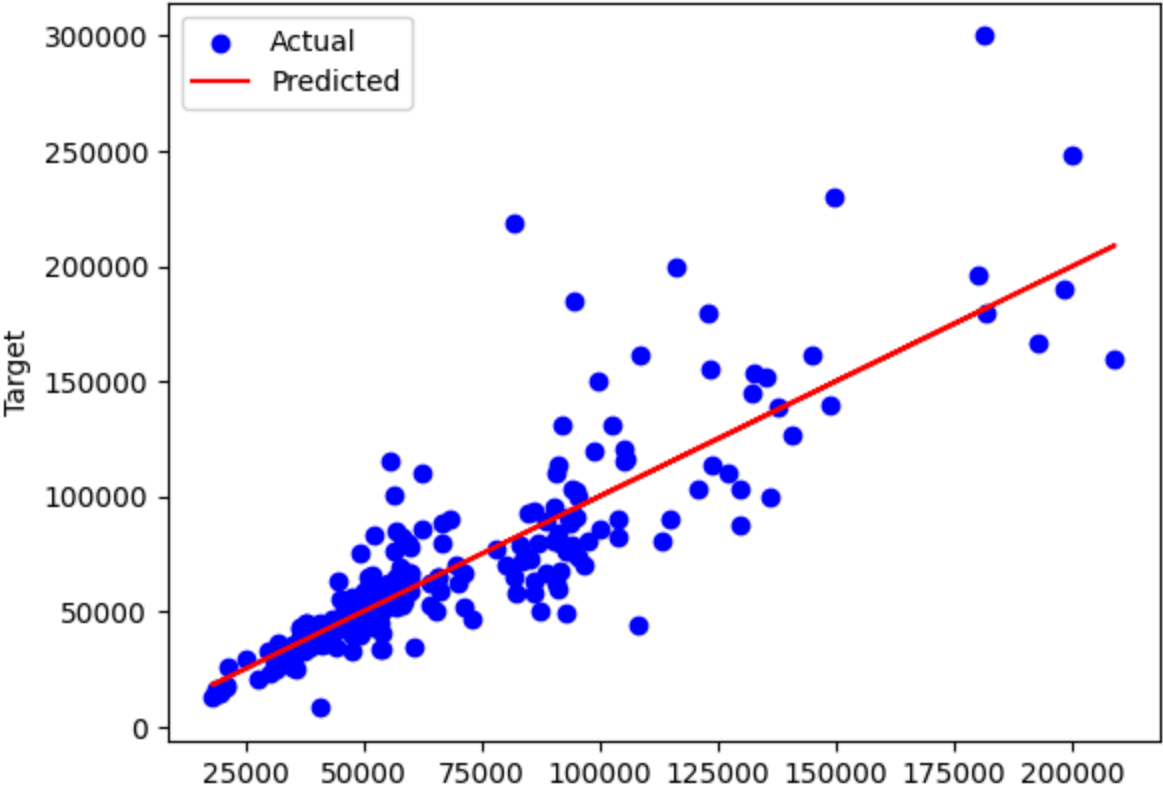
        y_pred = model.predict(X_test)
```

```
In [ ]: mse = mean_squared_error(y_test, y_pred)
        rmse = np.sqrt(mse)
        r2 = r2_score(y_test, y_pred)



        print("Root Mean Squared Error (RMSE):", rmse)
        print("R-squared Score (R^2):", r2)
```

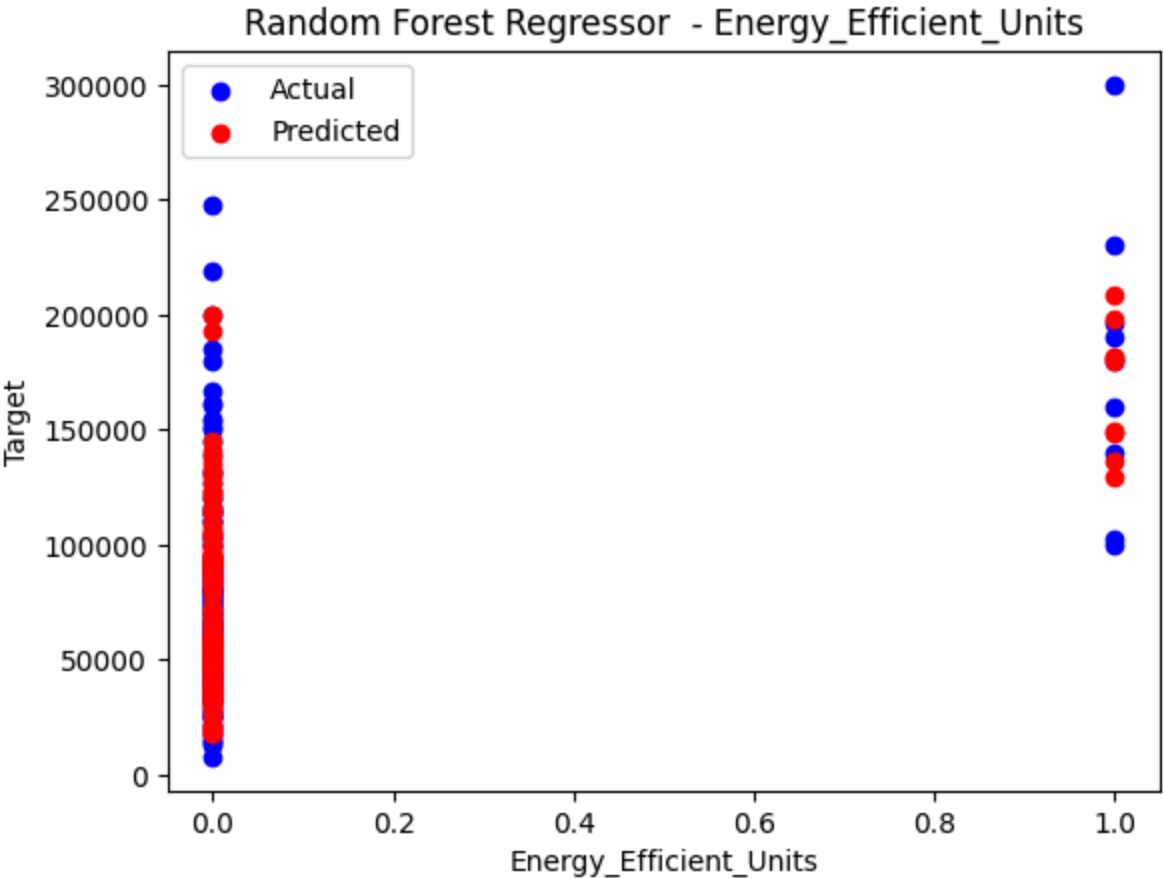

```
Root Mean Squared Error (RMSE): 24602.169847940197
R-squared Score (R^2): 0.720612527683302
```

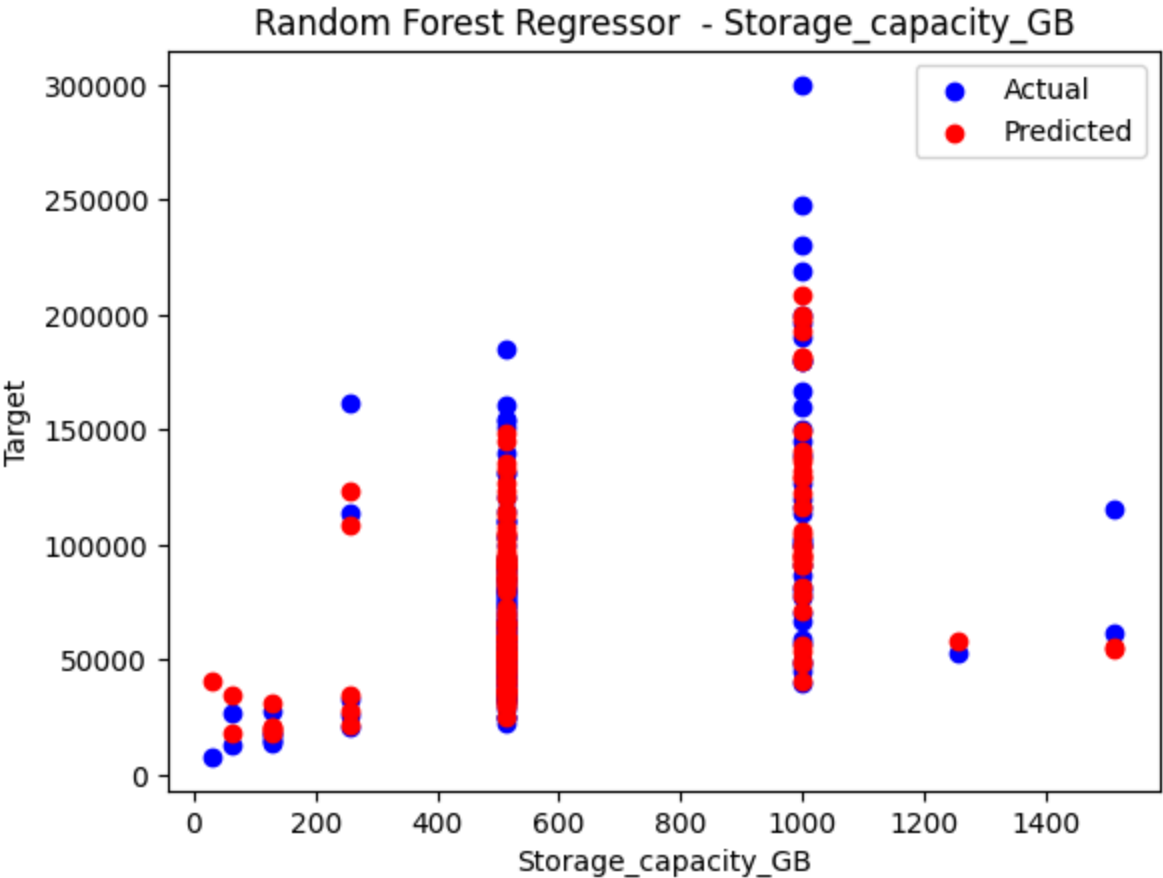
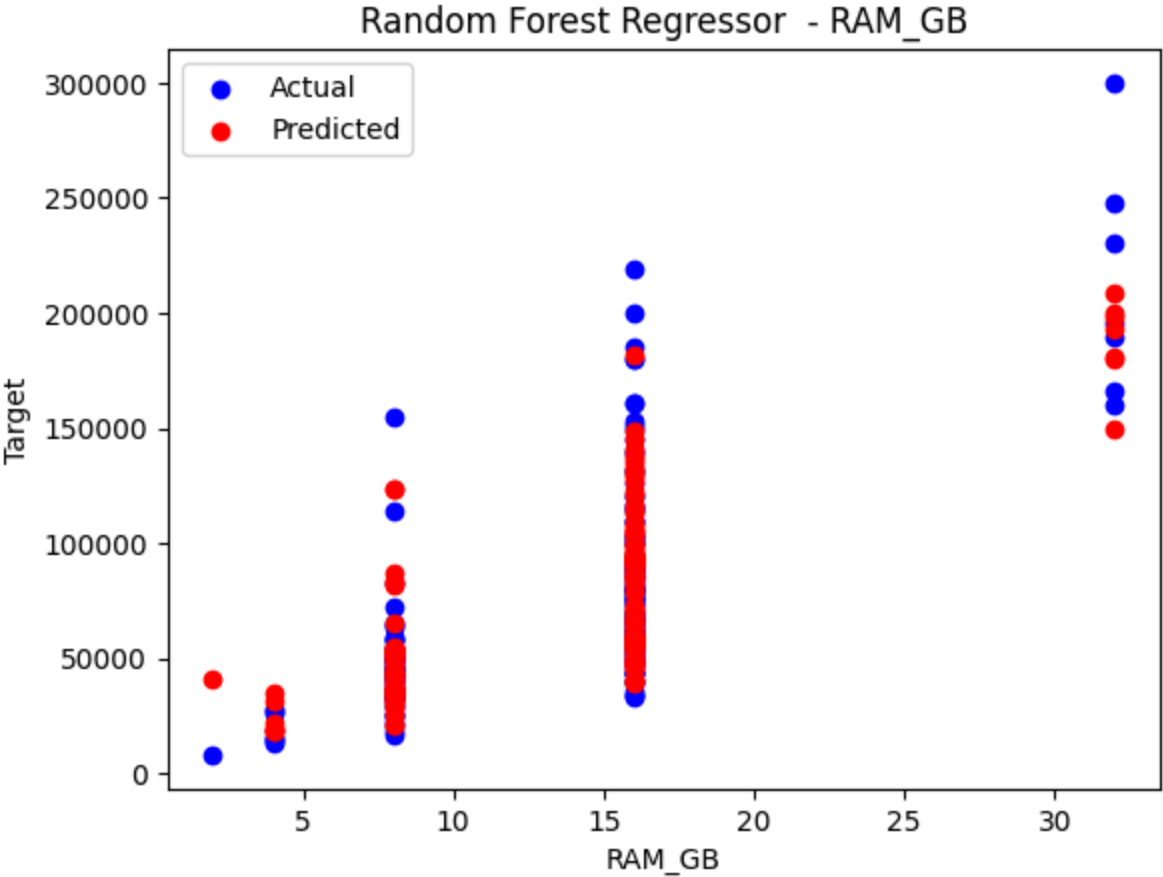
```
In [ ]: plt.scatter(y_pred, y_test, label='Actual', color='blue')
        plt.plot(y_pred, y_pred, label='Predicted', color='red')
        plt.ylabel('Target')
        plt.legend()
        plt.show()
```


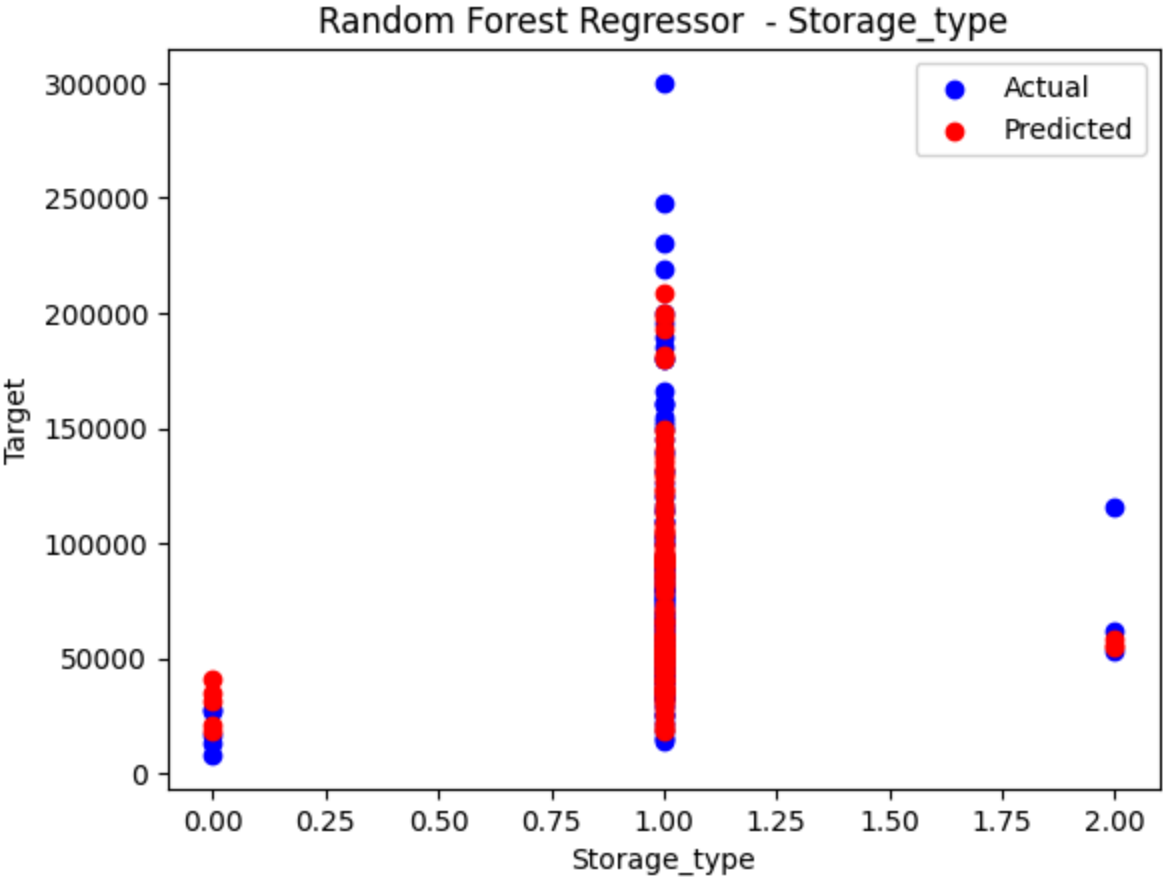


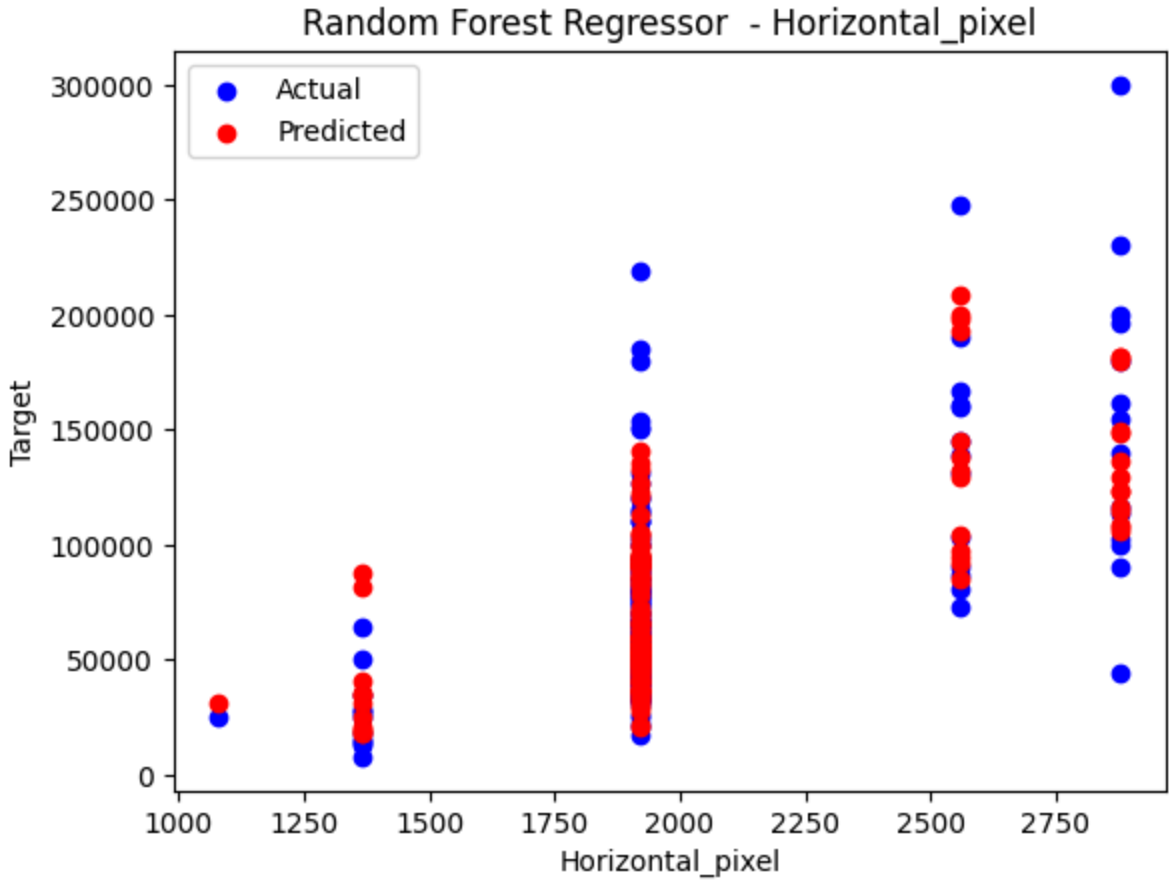
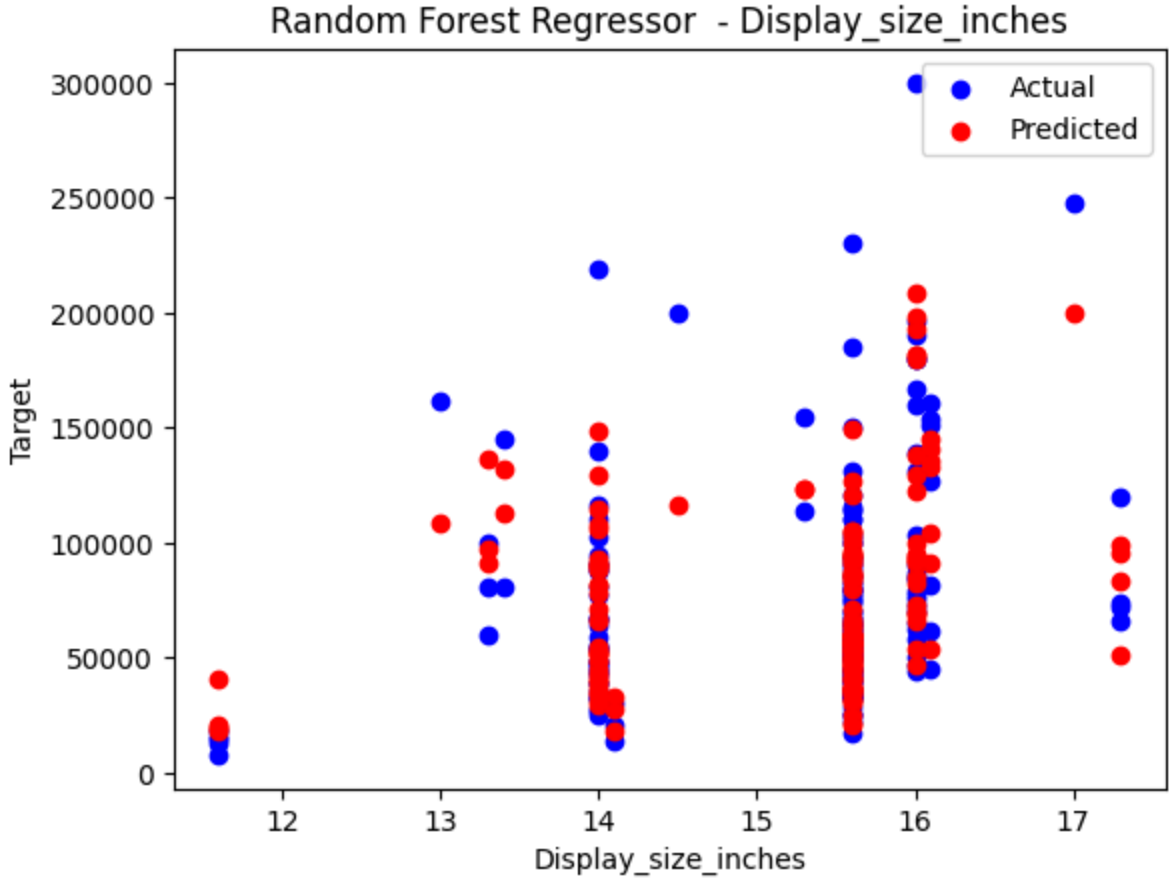
```
In [ ]: for column in X_test.columns:
    plt.scatter(X_test[column], y_test, label='Actual', color='blue')
    plt.scatter(X_test[column], y_pred, label='Predicted', color='red')
    plt.xlabel(column)
    plt.ylabel('Target')
    plt.title('Random Forest Regressor - {}'.format(column))
    plt.legend()
    plt.show()
```


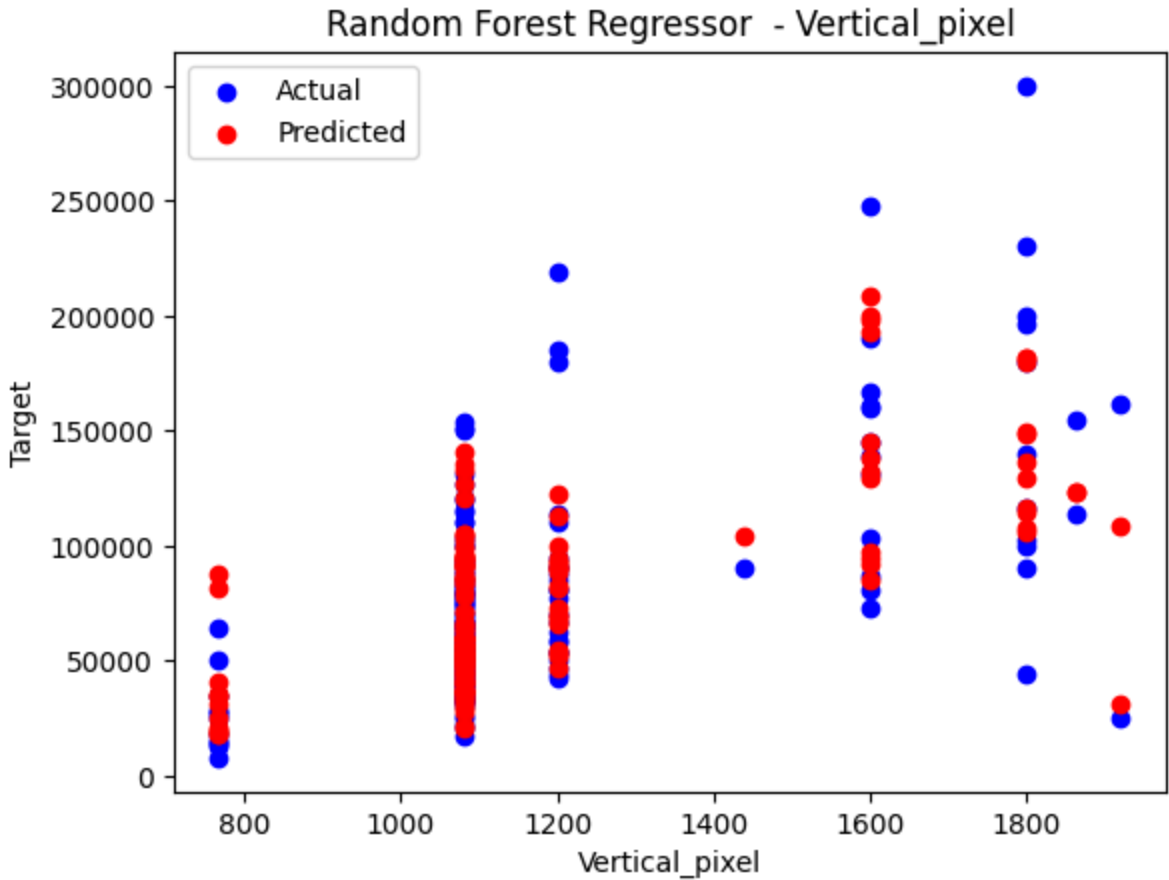



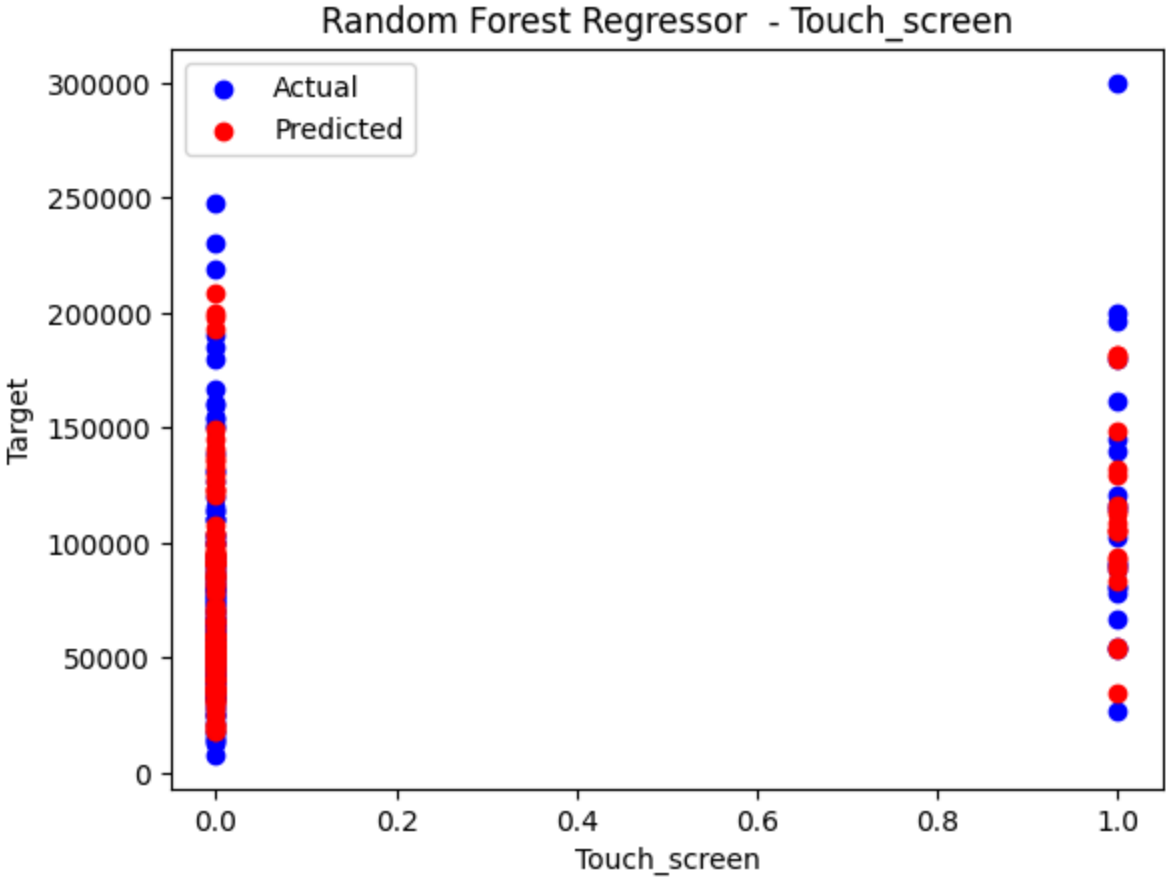












```
In [ ]: pd.DataFrame({'actual':y_test,'prediction':y_pred})
```

Out []:

	actual	prediction
322	102099	94847.773333
144	119990	98841.040000
518	33990	53811.467560
885	40300	53869.762833
91	57990	86152.061038
...
685	32990	35040.472083
654	51980	71193.458831
474	149990	99556.569167
558	82588	58266.870000
309	179990	181703.049333

198 rows × 2 columns

In []:

```
from sklearn.linear_model import LinearRegression

model = LinearRegression()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
```

In []:


```
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred)

print("Root Mean Squared Error (RMSE):", rmse)
print("R-squared Score (R^2):", r2)
```

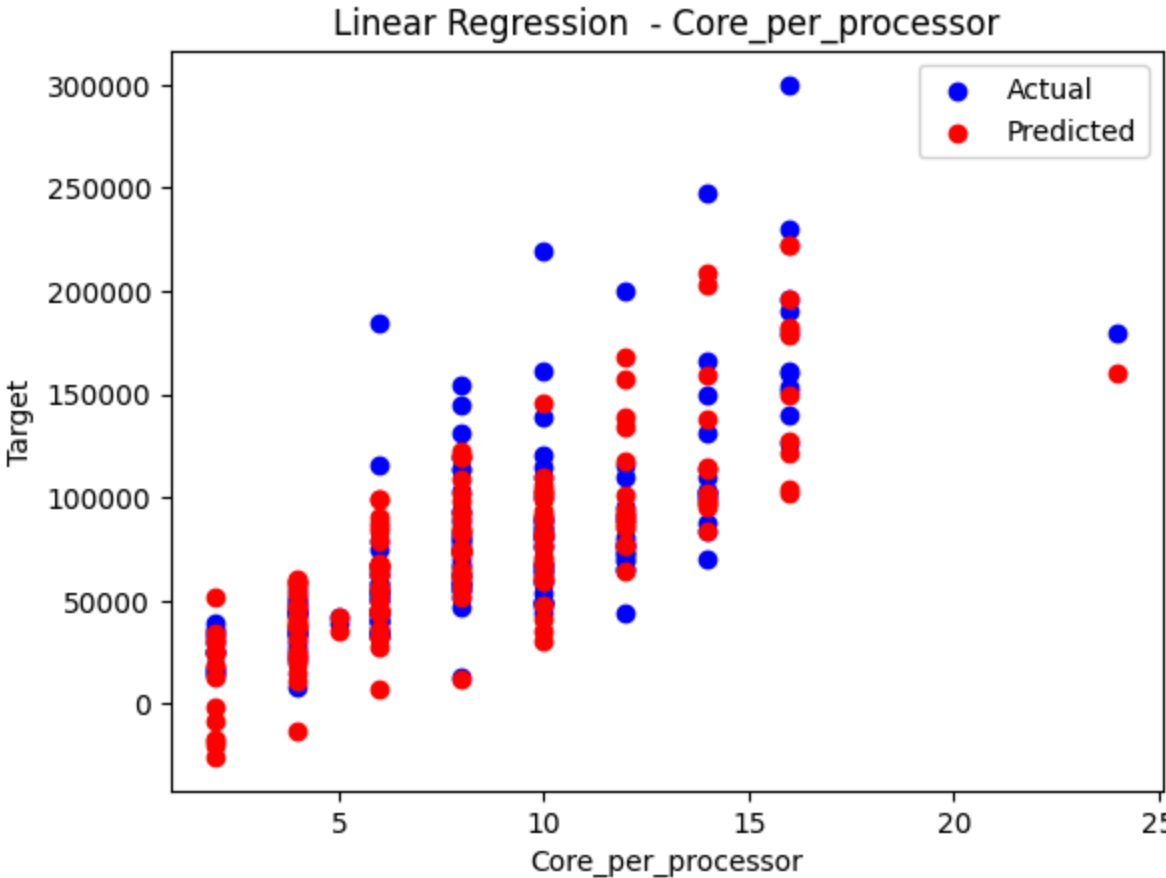

Root Mean Squared Error (RMSE): 23901.44784708819
R-squared Score (R^2): 0.7363009759228678



In []:



```
plt.scatter(y_pred, y_test, label='Actual', color='blue')
plt.plot(y_pred, y_pred, label='Predicted', color='red')
plt.ylabel('Target')
plt.legend()
plt.show()
```







```
In [ ]: for column in X_test.columns:
    plt.scatter(X_test[column], y_test, label='Actual', color='blue')
    plt.scatter(X_test[column], y_pred, label='Predicted', color='red')
    plt.xlabel(column)
    plt.ylabel('Target')
    plt.title('Linear Regression - {}'.format(column))
    plt.legend()
    plt.show()
```



















```
In [ ]: pd.DataFrame({'actual':y_test,'prediction':y_pred})
```

Out []:

	actual	prediction
322	102099	113720.801968
144	119990	89910.135318
518	33990	45315.910473
885	40300	54993.952549
91	57990	75245.653761
...
685	32990	36015.110443
654	51980	73213.739844
474	149990	114669.028463
558	82588	73620.122627
309	179990	178778.816933

198 rows × 2 columns

In []:

```
from sklearn.neural_network import MLPRegressor

# Assuming X_train, X_test, y_train, y_test are your training and testing data
model = MLPRegressor(hidden_layer_sizes=(100, 50), max_iter=5000)
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
```

In []:


```
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred)

print("Root Mean Squared Error (RMSE):", rmse)
print("R-squared Score (R^2):", r2)
```



Root Mean Squared Error (RMSE): 23572.58488186279
R-squared Score (R^2): 0.7435075888573958



In []:



```
plt.scatter(y_pred, y_test, label='Actual', color='blue')
plt.plot(y_pred, y_pred, label='Predicted', color='red')
plt.ylabel('Target')
plt.legend()
plt.show()
```







```
In [ ]: for column in X_test.columns:
plt.scatter(X_test[column], y_test, label='Actual', color='blue')
plt.scatter(X_test[column], y_pred, label='Predicted', color='red')
plt.xlabel(column)
plt.ylabel('Target')
plt.title('MLPRegressor - {}'.format(column))
plt.legend()
plt.show()
```



















```
In [ ]: pd.DataFrame({'actual':y_test,'prediction':y_pred})
```

Out[]:

	actual	prediction
322	102099	107471.520114
144	119990	94825.959023
518	33990	43908.637932
885	40300	60743.945626
91	57990	71267.180675
...
685	32990	43291.334124
654	51980	69402.273444
474	149990	108565.545032
558	82588	69775.254890
309	179990	173184.490433

198 rows × 2 columns